



FAKULTÄT FÜR INFORMATIK

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis

A Framework for Distributed Systems Based on The Actor Programming Model
and Dart language, Which Unifies Applications Across Devices, Clients and
Servers, and Supports Features for Hot Deployment and Migration of Actors

Sushil Man Shilpakar





FAKULTÄT FÜR INFORMATIK

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis

A Framework for Distributed Systems Based on The Actor Programming Model and Dart language, Which Unifies Applications Across Devices, Clients and Servers, and Supports Features for Hot Deployment and Migration of Actors

Ein Framework für verteilte Systeme auf der Basis des Actor-Programmiersmodells und der Dart-Programmiersprache, die Anwendung in Endgeräten, Clients und Servern erlaubt und die Möglichkeit für Hot Deployment und Migration der Actors bietet

Author:	Sushil Man Shilpakar
Supervisor:	Prof. Hans Arno Jacobsen
Advisor:	Richard Billeci
Submission Date:	December 15, 2014



I assure the single handed composition of this master's thesis only supported by declared resources.

Munich, December 15, 2014

Sushil Man Shilpakar

Acknowledgments

I would like to thank my supervisor Professor Hans Arno Jacobsen to providing me this wonderful opportunity to work in the thesis topic of my choice.

I will forever be thankful to my thesis advisor, Mr. Richard Billeci for introducing me to the Actor Programming model, helping me make useful decisions and providing me guidance throughout the project.

I would also like to thank Mr. Andreas Bucksteeg from Hybris GmbH for providing me access to Amazon EC2 servers during testing and benchmarking period of the project.

Abstract

The purpose of this thesis is to construct a framework that would allow developers to create concurrent, scalable and fault-tolerant applications with high availability in the Dart language.

The *isolate* of the Dart language is an interesting entity inspired by the actor model. Their nature of having no shared access to memory and relying on messages for communication makes them asynchronous and decoupled in nature. Nevertheless, their limitation of being able to be spawned only in a local Dart virtual machine restricts them from being distributed. The DDE (Dart Dart Everywhere) framework built as a result of this thesis uses the advantages of the Dart language enhancing it with libraries for remote management, hot deployment of code and distributed execution. The framework provides its own implementation of actors and allows developers to build applications that are inherently distributed in nature.

Contents

Acknowledgments	iii
Abstract	iv
1 Introduction	1
2 Literature Review	3
2.1 Message Passing Paradigm	3
2.2 Actor Programming Model	3
2.2.1 Error Handling in Actor Model	5
2.2.2 Differences from thread-based programming model	5
2.2.3 Scaling up actor based systems	6
2.3 Erlang	6
2.3.1 “Let it Crash” Philosophy	7
2.4 Scala Actors	7
2.5 Akka Toolkit	8
2.5.1 Actor Model in Akka	8
2.5.2 Actor System	9
2.5.3 Message Passing in Akka	10
2.5.4 Shared mutable state	11
2.5.5 Actor Supervision and Monitoring	11
2.5.6 Routers	12
2.5.7 Remote Actors in Akka	12
2.5.8 Clustering in Akka	14
2.6 The Dart Language	14
2.6.1 Overview	14
2.6.2 Advantages of Dart	15
2.6.3 Dart and JavaScript	15
2.6.4 Asynchronous Programming in Dart	15
2.6.5 Isolates	15
2.7 RabbitMQ - A Message Broker System	18
2.7.1 Message Queues in RabbitMQ	20

2.7.2	RabbitMQ and prefetch-count	20
2.8	STOMP	21
2.8.1	Protocol Overview	21
2.8.2	STOMP Library for Dart	22
2.9	WebSocket	22
2.9.1	Security	22
2.9.2	Establishing a Connection	22
3	System Design	24
3.1	Core Design Decisions	24
3.2	The Framework	25
3.2.1	Isolate System	26
3.2.2	The Registry	33
3.2.3	Message Queuing System (MQS)	36
3.2.4	Activator	38
3.3	Some Key Features	39
3.3.1	Hot Deployment of Isolates and Isolate Systems	39
3.3.2	Migration of Isolates and Isolate Systems	39
3.3.3	Remote Isolates	40
3.4	Typical Message Flow in the System	40
3.4.1	Sample message formats	41
3.4.2	Some Implementation Overview	43
3.4.3	Clustering	45
3.5	Dart Libraries Used in Construction	45
3.6	A Sample Implementation of Worker Using The Framework	46
4	Results	47
4.1	Applications based on the DDE framework	47
4.1.1	Producer/Consumer Program	47
4.1.2	Requester/Supplier Program	49
4.1.3	System Setup for Testing	51
4.1.4	Observations from DDE based applications	51
4.2	Benchmarks	52
4.2.1	Prefetch Count	52
4.2.2	Message Size	54
4.2.3	Number of Message Queuing Systems	55
4.2.4	Production Throughput of Isolates	57
4.2.5	Simultaneous Production and Consumption	58
4.2.6	Throughput of Request-Reply	60

4.2.7	Round Trip Time in Request-Reply	61
5	Discussions	62
5.1	Performance Findings	62
5.1.1	Effect of Prefetch Count	62
5.1.2	Effect of Message Size	62
5.1.3	Effect of Scaling out the Message Queuing System	63
5.1.4	Comparision of production throughputs	63
5.1.5	Effect of Simultaneous Production and Consumption	64
5.1.6	Throughput of Request-Reply	64
5.1.7	Round Trip Time of Request-Reply	64
5.2	The DDE Framework	65
5.3	Problems/Issues	65
5.3.1	Dart Induced Issues	65
5.3.2	STOMP Library for Dart	66
6	Conclusion	67
7	Future Directions	68
	Acronyms	69
	List of Figures	70
	List of Tables	71
	Bibliography	72

1 Introduction

To scale up and handle a high number of requests, applications are deployed in distributed environments. The transfer speed and demand for data have been increasing rapidly. Performance expectations from web applications are higher than ever. Performance and scalability concerns are met by using high end multi-core servers and concurrent processing via multi-threading.

Enterprises want their applications to be available 100% of the time because a few seconds of application downtime could cause a huge loss to revenue and customer satisfaction.

Ideally an application should be able to scale up or down without compromising its availability. The possibility of dynamically adjusting scalability allows an application to grow when demand increases and shrink to minimal resources when demand is low.

Deploying applications to distributed environments is a good solution for scaling up, but most applications fail to utilize all available physical resources because of their underlying design.

The actor programming model [section 2.2] allows developers to build highly concurrent, distributed and fault tolerant event-driven applications. These characteristics are provided by actors which use asynchronous message passing [section 2.1]. Actor implementations were first created in 1980s, but only recently have become mainstream.

The work presented in this thesis focuses on creating a framework for the Dart language [section 2.6] that allows developers to create asynchronous, concurrent, scalable and distributed applications. The concept of scalability in the framework is based on the actor-like nature of isolates [subsection 2.6.5], which communicate solely by message passing. The framework takes advantage of isolates and extends its functionality so that they can be deployed into distributed systems.

The framework is intended to take advantage of the fact that the Dart virtual machine can be run in browser as well as in server. This advantage opens up the possibility of creating a fully distributed application in which isolates may run everywhere: in servers, desktop browsers and even in mobile browsers.

Erlang [section 2.3] and Scala [section 2.4] are popular programming languages which have their own implementation of actor programming. The akka toolkit [section 2.5] provides a foundation for actor programming in both Java and Scala. The DDE framework presented in this thesis draws on ideas from the actor system [subsection 2.5.2]

and router [subsection 2.5.6] implementations of akka to provide a comparable isolate system [subsection 3.2.1] and routers [section 3.2.1] for Dart.

As applications of the DDE framework are based on message passing, the framework takes advantages of features offered by the message broker system – RabbitMQ [section 2.7]. RabbitMQ provides message persistence and a decoupled nature to applications built on the DDE framework.

To provide a fully distributed nature to applications the DDE framework allows developers to deploy components of their application (Worker isolate [section 3.2.1]) to remote nodes, which is achieved by serializing messages through WebSockets [section 2.9]. The DDE framework itself is based on asynchronous message passing. It has several decoupled components which communicate with each other via WebSockets [section 2.9].

The result of this thesis is a framework based on the Dart language that enforces actor based programming in Dart and makes applications scalable, highly available, with support for deploying and terminating isolates at runtime without interrupting the whole system.

Chapter 2 consists of the introduction and references of the concepts behind the DDE framework. It also discusses the technologies that were used in the framework. Chapter 3 presents the implementation ideas and proof of concept of Dart Everywhere. Chapter 4 shows sample applications built using this framework and presents performance benchmarks and evaluations of applications ran in a distributed environment in Amazon EC2. The deductions made from the benchmarks and results are discussed in chapter 5.

2 Literature Review

2.1 Message Passing Paradigm

Message passing is sending a message from one process, component or actor to another. When a message is received by a recipient, it chooses further processing based on the pattern or content of the message.

Message passing is the loosest type of coupling. The components of a system are not dependent on each other; instead they use a public interface to exchange parameterless messages[Cell11]. Hence, systems that implement message passing paradigm are easily scalable and efficient. Such systems are easy to replicate and easy to make fault-tolerant. [Arm10].

Systems which communicate solely by message passing do not have a shared state. Such systems can be easily divided into isolated components. This makes the overall architecture easy to understand and simplifies the isolation of problems within the system [Arm10].

Message passing can either be synchronous or asynchronous. Synchronous message passing is the communication between two components where both the sender and intended receiver of a message are ready to communicate. But, as the sender of a message has to wait for a response from the receiver, the sender is usually blocked until it gets a reply. Whereas, in asynchronous message passing, the sender simply sends a message and continues to do other tasks. The receiver does not have to be ready to accept the message when the sender sends it [Agh85]. Thus, asynchronous message passing is non-blocking in nature.

2.2 Actor Programming Model

Actor programming model is a programming paradigm designed for concurrent computation. The concept of an actor was originally introduced by Carl Hewitt[AH85]. He, along with Agha[Agh85] have been involved in development of the actor theory as well as its implementation.

An actor is a fundamental unit of computation. It is neither an operating system process nor a thread, but a lightweight process. It embodies three essential things:

- Processing
- Storage
- Communication

Actors have addresses and there is a many-to-many relationship between actors and addresses.

Actors communicate with each other in a non-blocking way by asynchronous message passing, which removes the need of explicit locks. An actor can send message to actors in same system or another system. An actor can also send message to itself, which is how a recursion is achieved. An actor can send a message to target actor only if it has the address of target actor. Agha ([Agh85], p35) lists three ways in which an actor, upon accepting a message, can know the address of the target actor:

- the target was known to the actor a before it accepted the message
- the target became known when the message was accepted because it was contained in the message
- the target is the new actor created as a result of accepting the message

To buffer incoming messages, each actor has a mailbox. A mailbox is a queue of messages that have been sent by other actors or processes and not yet consumed, where mailbox is also an actor. According to Hewitt, the order in which the messages are delivered is non-deterministic [Nin].

After receiving a message, an actor may perform following actions [AH85]:

- Create other actors
- Send messages to itself, other known actors or reply to the actor who sent the message
- Designate how it is going to handle next message, i.e. Specify a replacement behavior

Actors do not have a shared mutable state. All mutable state is private to the actor and all shared state is immutable. Actors communicate with each other by asynchronous message passing which is also immutable. Each actor processes only one message at a time, and unless it is a broadcast message, a message is not processed multiple times.

An actor exists in a system. An actor system is a group of actors working together in certain hierarchy.

In the actor model, concurrency is inherent because of the way it is designed. Also, there is no guarantee that message sent to an actor will arrive sequentially [Nin]. Nevertheless, Akka framework [section 2.5] guarantees the order of delivery of messages between two actors, provided that there are no intermediary actors in between [subsection 2.5.3].

Several programming languages like Act 1, 2 and 3, Acttalk etc. were created when actor system was newly introduced by Hewitt and Agha [Agh85; AH85].

2.2.1 Error Handling in Actor Model

Exception handling in actor model is based on idea of embracing failures. The idea of embracing failures is also known as “let it crash” paradigm [subsection 2.3.1]. As the actors do not have shared state, this allows individual actor to fail without causing disruptions in the system. Since, an actor in an actor system is typically organized in a hierarchical structure, the actor which created a child actor can be used for supervising it. The idea of supervision in a hierarchical actor system helps to make the actor system fault tolerant [Erb12]. When an actor throws an exception the supervisor can respond to the exception in different ways. In Akka [section 2.5], the supervising actor usually reacts by either simply ignoring the exception and letting the actor continue, by restarting the actor or by escalating the error to its supervisor.

The “let it crash” style of programming is a non-defensive programming, which is implemented successfully in Erlang [section 2.3].

2.2.2 Differences from thread-based programming model

Thread-based Concurrency

In thread-based programming languages, the control flow of a program is divided into several threads for concurrency. The threads operate simultaneously and the control can switch from one thread to another non-deterministically. When two or more threads have shared memory, concurrent modifications and accesses of data might result in undesired behavior of the system, known as ‘race condition’. To prevent this type of situation, such programming languages use locks. The locks let only a single thread at a time run sequentially for the section of a program code [Sofa].

Generally, thread-based programming models are easy to understand and implement. But, the resulting program behavior is difficult to understand because of implicit context switches and release of locks, which may lead to a deadlock situation [Sofa].

Actors based Concurrency

Concurrency in actor based programming languages are inherent because of asynchronous message-passing, pipelining, and the dynamic creation of actors. The concurrency in actors through pipelining is only constrained by the logical limits and the available hardware resources [Agh85]. The actors may carry out their activities in parallel as each actor resides in a completely separated space from the other actors, they are connected only via messages.

Actor based programming liberates the programmer from delving into coding details about the parallelism and threads [Agh85].

The ‘race condition’, discussed in section 2.2.2, do not arise in actor based programming as actors do not have a shared state and they don’t need locks.

2.2.3 Scaling up actor based systems

As actor based systems are highly concurrent [section 2.2.2], it is easy to scale actor based systems from a single core system to several of them across multiple data centers around the globe. An ideal case would be to just let a new node join the actor system and let it run some more actors during runtime, without taking the running system offline. The actors themselves do not need to know the physical location of other actors as they simply exchange messages based on the logical addresses. This makes them easy to scale up by adding more actors to the same machine and scale out by adding more nodes, running actors, to a distributed system.

2.3 Erlang

Erlang is the first popular programming language based on the actor model [Vin07]. It was developed by Joe Armstrong in 1986 at the Ericsson Computer Science Laboratory and was made open-source in 1998. It was chiefly used in telephony applications as it was built to solve the problems of availability as well as scalability that existed in such applications [Arm07].

Erlang ‘processes,’ which are essentially user-space threads rather than Unix processes or kernel threads, communicate only via message passing.

Erlang was designed for the writing applications that require high availability [Arm07]. It uses concurrent processes, which have no shared memory and communicate only via asynchronous message passing. This idea is similar to the actors model proposed by Agha and Hewitt [section 2.2]. Thus, programs written in Erlang are concurrent, distributed, fault tolerant and thread safe. The concurrency is built into the language itself, not the the operating system [Arm10].

When an application developed in Erlang is deployed in a multicore computer, it automatically takes advantage of those multiple cores. The Erlang processes distribute over the cores. So, programmers do not have to worry about threads [Arm10]. In Erlang, new changes in an application can be added to the system without taking it offline. This improves the availability of whole system. Thus, it simplifies the construction of software for implementing non-stop systems [Arm07].

Error handling in Erlang is different from most of the other programming languages. The error handling is based on a “let it crash” philosophy [subsection 2.3.1] which is a non-defensive style of programming [Arm07; Arm10].

2.3.1 “Let it Crash” Philosophy

The core idea in “Let it Crash” philosophy is to let the failing processes crash and make other processes, which observe this process, detect the crashes and fix them [Arm10]. This idea is in sharp contrast to other programming languages, where programmers implement exception handlers and prevent a process from getting terminated.

The proponents of the “let it crash” philosophy argue that it leads to more clear and compact code. [Arm10].

2.4 Scala Actors

Scala is a statically typed programming language which integrates functional as well as object-oriented programming [OR14]. Scala has its own library for actor programming.

In Scala, templates for actors with user-defined behavior are normal class definitions which extend the predefined Actor class.

The scala actors library provide concurrent programming model based on actors [section 2.2]. The actors in scala are fully inter-operable with the ordinary virtual machine threads [HO09]. Scala actors are lightweight and support around 240 times more actors to run simultaneously compared to virtual machine threads [HO09].

Both synchronous and asynchronous message passing can be used in scala actors. Synchronous message passing is implemented by exchanging several asynchronous messages. Actors in scala can also communicate using ‘futures’. When a future is used the requests are handled asynchronously and the sender immediately gets a representation of the future which allows sender to wait for the reply [Hal].

The Scala Actors library is now deprecated and will be removed in future Scala releases¹. The deprecation is in favor of the use of Akka actors [section 2.5].

¹Scala actors are deprecated in version 2.10 [HT]

2.5 Akka Toolkit

Akka is a toolkit and runtime for building highly concurrent, distributed and resilient message-drive applications on the JVM [Incb]. Akka uses lightweight actors for concurrency. The actors in Akka are based on Hewitt and Agha's model [Agh85; AH85] of actor programming. Akka provides a well defined API for developers to create large concurrent systems, which allows for easy scaling out. Akka is available for Scala as well as Java.

2.5.1 Actor Model in Akka

In Akka, actors are objects which encapsulate state and behavior. Similar to Hewitt and Agha's actor model [section 2.2], the akka actors communicate only by messages which are placed into the recipient's mailbox. Actors, in akka, are objects. Therefore, Akka actors enforce a stringent form of object-oriented programming [Inca].

The 'ActorRef' in akka, represents an actor. It holds a reference to an actor. Its purpose is to facilitate the message sending to the actor it refers. An actor can also refer to itself using *self* [Incc].

The Listing 2.1², is a simple example of how actor programming in Akka is realized. In this example, the 'GreetingActor' which extends 'UntypedActor' must implement the *onReceive()* method. The *onReceive()* method is invoked for each message received by this actor. After receiving a message, the actor performs pattern matching to decide which code to execute. As in this example, if the message is a type of 'Greeting' object, it is concatenated with the string "Hello" otherwise, it is simply ignored.

Listing 2.1: A simple example of actor programming in akka [Incb]

```
public class Greeting implements Serializable {
    public final String who;
    public Greeting(String who) { this.who = who; }
}

public class GreetingActor extends UntypedActor {
    LoggingAdapter log = Logging.getLogger(getContext().system(), this);

    public void onReceive(Object message) throws Exception {
        if (message instanceof Greeting)
            log.info("Hello " + ((Greeting) message).who);
    }
}
```

²Source: T. Inc. *Akka Toolkit*. <http://akka.io>. Last Accessed: 2014-11-19


```
ActorSystem system = ActorSystem.create("MySystem");
ActorRef greeter = system.actorOf(Props.create(GreetingActor.class), "greeter");
greeter.tell(new Greeting("Charlie Parker"), ActorRef.noSender());
```

2.5.2 Actor System

An actor system arranges actors in a hierarchical structure. The hierarchical structure of actors are formed when an actor starts up child actors to handle certain tasks. The actor not only creates and assigns tasks to the child actors but also supervises them [subsection 2.5.5]. Thus it is implicit that there can be only one supervisor of an actor.

By splitting up the tasks, the tasks become clear and structured. Furthermore, the resulting actors also becomes simplified and specialized in terms of which messages it should process and how it should react normally. It also becomes easy to handle the failures. If an actor cannot cope with a failure, it is escalated up in the hierarchy to its supervisor. The escalation is repeated by actors in every higher level hierarchy unless it reaches up to the actor which can handle the failure [Incc].

An actor instance in akka takes up roughly 300 bytes of memory, because of which it is possible to spawn millions of them in one actor system. Akka can manage very large and distributed systems; it can take care of the order in which messages are processed even in large systems [Incc].

An example of the arrangement of actors in hierarchical structure is shown in Figure 2.1. The `\` is known as the “root guardian” and `\user` is known as the “user guardian”. Any actor created by a user falls under the user guardian. For instance, in the Figure 2.1, *actorA* and *actorB* are actors created by user, hence their supervisor is `\user`. Again, the actors *actor1* and *actor2* are the child actors of *actorB*. The address of an actor in akka is arranged like the filesystem hierarchy. Hence, the supervision hierarchy as well as the path leading to the actor is comprehensible from its address.




Figure 2.1: Hierarchy in actor system [Incc]

2.5.3 Message Passing in Akka

In Akka, every actor has an event driven message inbox known as the ‘mailbox’. The mailbox buffers all incoming messages until they are processed. When a message is sent from one actor to another, the reference to the sender (an ActorRef), is automatically added to the message by default. Thus, during the message processing the recipient actor has a reference to the sender actor through the *sender* method [Incc].

Akka guarantees the order of direct message delivery between two actors. It is only applicable when the mailbox implementation in receiver is FIFO mailbox and the communication takes place only between the two actors without the involvement of intermediary actors. For instance, if an actor A1 sends messages M1, M2 and M3 to actor A2. The message M1 will be delivered before M2 and the message M2 before M3. Meanwhile, if another actor A3 sends messages M4, M5 and M6 to A2 at the same time, the order of message delivery of M4, M5 and M6 will also be sequential. Nevertheless,

there is no guarantee that the messages sent by A1 is delivered before the messages from A3; even if all of the messages from A1 were sent before A3 started sending messages to A2. Provided that the message sending from A1 and A3 are independent of each other [Incc].

The Akka's documentation for Java, the rules mentioned for message sending are: [Incc]

- at-most-once delivery, i.e. no guaranteed delivery
- message ordering per sender–receiver pair

Here, 'at-most-once delivery' means that a message is either not delivered at all (i.e. lost) or delivered only once. Thus, there is no duplicate delivery of a message.

2.5.4 Shared mutable state

Since Akka runs on top of the Java Virtual Machine(JVM), there are some pitfalls that the programmer should be aware enough to avoid, which akka cannot enforce. The messages used to communicate between the actors should be immutable. If it is made mutable and the reference of it is sent by the sender to an actor which resides on same JVM, bugs such as race conditions and even some incomprehensible bugs might appear. Also, the sender method must not be closed over if the block of code could be run in another thread; for instance, sending (replying) a message to the sender actor inside a 'Future' block. In such cases, the sender reference must be captured into a local variable first. As the reference and behavior of sender may change over time [Incc].

2.5.5 Actor Supervision and Monitoring

The supervising actor in akka delegates tasks to its subordinates. It is also responsible for monitoring them for failures. When a child actor detects a failure, it suspends itself as well as its child actors and sends a message to its supervisor about the failure. According to Akka Java documentation[Incc], upon receiving the failure, the supervisor may opt to perform one of the four choices:

1. Resume the subordinate, keeping its accumulated internal state
2. Restart the subordinate, clearing out its accumulated internal state
3. Stop the subordinate permanently
4. Escalate the failure, thereby failing itself

When an actor is restarted by the supervising actor, the message that the actor was processing during the time of failure is lost and is not processed again. However, the messages that were in the mailbox of the actor remains safe and new actor resumes processing the other messages in the mailbox [Incc].

2.5.6 Routers

Routers are specialized actors that act as a load balancer for the actors that it supervises. Akka provides several built-in routing techniques: [Incc]

- Round Robin Router
- Random Router
- Smallest Mailbox Router
- Broadcast Router
- Scatter Gather First Completed Router
- Tail Chopping Router
- Consistent Hashing Router

It is also possible to use custom implementation of a router by extending *RoutingLogic* class from Akka's routing library.

2.5.7 Remote Actors in Akka

Actors in akka are location transparent; they reside in a logical hierarchy of an actor system through which the physical location of an actor in the network can be determined. The location transparency allows the akka applications to be developed locally but deployed into distributed systems simply by changing configurations [Incc].

The peer-to-peer communication between actor systems is the foundation for Akka Clustering. Akka's Documentation for Java [Incc] lists the two design decisions for remoting:

1. Communication between involved systems is symmetric: if system A can connect to system B then system B must also be able to connect to system A independently.
2. The role of the communicating systems are symmetric in regards to connection patterns: there is no system that only accepts connections, and there is no system that only initiates connections.

The Listing 2.2 is an example taken from website of Akka³ which is an example for configuration and code for deploying actors in remote nodes.

Listing 2.2: A sample Akka Configuration and Code for Remote Actors [Inc]

```
// -----
// config on all machines
akka {
  actor {
    provider = akka.remote.RemoteActorRefProvider
    deployment {
      /greeter {
        remote = akka.tcp://MySystem@machine1:2552
      }
    }
  }
}

// -----
// define the greeting actor and the greeting message
public class Greeting implements Serializable {
  public final String who;
  public Greeting(String who) { this.who = who; }
}

public class GreetingActor extends UntypedActor {
  LoggingAdapter log = Logging.getLogger(getContext().system(), this);

  public void onReceive(Object message) throws Exception {
    if (message instanceof Greeting)
      log.info("Hello " + ((Greeting) message).who);
  }
}

// -----
// on machine 1: empty system, target for deployment from machine 2
ActorSystem system = ActorSystem.create("MySystem");

// -----
// on machine 2: Remote Deployment - deploying on machine1
ActorSystem system = ActorSystem.create("MySystem");
ActorRef greeter = system.actorOf(Props.create(GreetingActor.class), "greeter");

// -----
// on machine 3: Remote Lookup (logical home of "greeter" is machine2, remote
// deployment is transparent)
```

³T. Inc. Akka Toolkit. <http://akka.io>. Last Accessed: 2014-11-19

```
ActorSystem system = ActorSystem.create("MySystem");
ActorSelection greeter =
    system.actorSelection("akka.tcp://MySystem@machine2:2552/user/greeter");
greeter.tell(new Greeting("Sonny Rollins"), ActorRef.noSender());
```

2.5.8 Clustering in Akka

Akka does not have a central server. Instead, it uses peer-to-peer based gossip protocol to form a cluster. Thus, there is no single point of failure or single point of bottleneck in the system. But, adding a new member in a distributed system requires significant amount of time because of the nature of the gossip protocol. In a benchmark [Nor] performed by Patrik Nordwall⁴, adding nodes to a cluster of at least 1500 nodes took around 15 to 20 seconds in average.

2.6 The Dart Language

2.6.1 Overview

Dart is an open-source, class-based, single-inheritance, pure object-oriented programming language developed by Google [ECM14]. Dart language is inspired by Smalltalk, Strongtalk, Erlang, C# and JavaScript [Lad].

Dart code is not compiled before running; it is read and executed directly from the source code by the Dart Virtual Machine (VM). Dart provides a homogeneous system that encompass both client as well as server as the Dart VM can be embedded in browsers. A version of Chromium – ‘Dartium’ already has Dart VM built into it [Lad].

Programs in Dart are optionally typed. They can be executed in two modes checked mode and production mode. In checked mode incorrect static type annotations produce compile time errors. Whereas, in production mode, type annotations are completely ignored [ECM14].

Furthermore, Dart allows developers to code in a uniform way for both server as well as client since Dart Virtual Machine (VM) can be embedded in browsers. A variant of Chromium browser — Dartium browser has an embedded Dart VM.

Dart has automatic garbage collecting system, which means the memory occupied by objects which are not in use and which do not have any reference are reclaimed periodically.

Some of its important features are:

- Easy to learn syntax

⁴Patrick Nordwall is a developer of Akka at Typesafe Inc.

- Compiles (Translates) to JavaScript
- Runs in client as well as on server
- Dart supports types, but it is optional
- Can scale from small script to large and complex applications
- Support safe concurrency with isolates
- Support of code sharing

2.6.2 Advantages of Dart

- Translates to JavaScript so that the code can be run in the web-browsers that do not have Dart VM yet
- Has been gaining popularity and adoption in recent years
- Optionally typed language

2.6.3 Dart and JavaScript

Codes written in Dart can be translated to JavaScript using a tool – ‘dart2js’. ‘Dart2js’ is bundled with the Dart SDK (Software Development Kit). As popular web browsers like Mozilla Firefox, Google Chrome, Safari do not have Dart Virtual Machine embedded, the ability to translate source code from Dart to JavaScript lets dart programs run in any modern browser without needing to manually port the source code to JavaScript.

2.6.4 Asynchronous Programming in Dart

Most programming languages use callback functions for asynchronous programming. Dart provides some additional alternatives along with callback functions – Future and Stream objects. A ‘Future’ is a promise for a result which will be returned after an arbitrary amount of time. A ‘Stream’ is a way to get a sequence of values, such as events, data from ports etc.

2.6.5 Isolates

Although Dart programs are single threaded, concurrency is supported via actor-like entities called isolates. An isolate has its own memory and own thread of control. Message passing is the sole way to communicate between isolates. No state is ever shared between isolates. Isolates are created by spawning [ECM14].

An isolate has its own heap memory different from the main isolate (the top level isolate). It is possible for the child isolate to throw exceptions and errors such as by exhausting its memory. If the exceptions are not handled properly, it forces the isolate to be shutdown.

In Dart, when an isolate is spawned, usually a ‘sending port’ is sent as the initial message (by the spawner), so that spawner and “spawnee” can communicate with each other. The “spawnee” can use the sending port sent by spawner to reply to the spawner.

An isolate can spawn another isolate which can further spawn other isolates and have control over them. Thus, the spawner can supervise the “spawnee”. The spawner can pause the “spawnee” or terminate it [Goo]⁵.

Modern web browsers, even on mobile platforms, run on multi-core CPUs. To take advantage of all those cores, developers traditionally use shared-memory threads running concurrently. However, shared-state concurrency is error prone and can lead to complicated code. Thus, instead of threads, all Dart code runs inside of isolates. Each isolate has its own memory heap, ensuring that no isolate’s state is accessible from any other isolate [Kat12].

Spawning an Isolate

There are two ways to spawn an isolate: using *Isolate.spawnUri()* or using *Isolate.spawn()*. The *Isolate.spawn()* uses top level function to spawn an isolate. The top level function may reside in the same class or may belong to another class. The *Isolate.spawnUri()* spawns an isolate using the source code of a file from a given location. The location can be a remote http/https URI or a path to source file in local disk. To spawn an isolate using *Isolate.spawnUri()*, the source file must have an entry point function *main()*. The newly spawned isolate shares the same code as the spawner isolate [Goo].

Communication Between Two Isolates

After an isolate is spawned, it is recommended to send its *SendPort* to the spawner isolate so that the “spawner” and “spawnee” can communicate via message passing. *SendPort* and *ReceivePort* are used by the isolates to communicate with each other.

The example in Listing 2.3 shows how an isolate is spawned using *spawnUri()* and how to perform basic communications between two isolates in Dart. As shown in this example, the message sending is performed after the *SendPorts* are exchanged between the isolates.

Listing 2.3: A simple example of isolate communication in dart

⁵Pausing and terminating an isolate is not available in Dart 1.7.2 or older versions


```
//sample.dart
import 'dart:isolate';

main(var args, SendPort sendPort) {
  ReceivePort receivePort = new ReceivePort();
  SendPort sendport;
  sendPort.send(receivePort.sendPort);

  receivePort.listen((var message) {
    if(message is SendPort) {
      sendPort = message;
    } else if(message is String) {
      print("Received: $message");
    } else {
      sendPort.send("Unknown Message");
    }
  });
}

//app.dart
import 'dart:isolate';

main() {
  ReceivePort receivePort = new ReceivePort();
  SendPort sendPort;
  Isolate.spawnUri(Uri.parse("sample.dart"),null,receivePort.sendPort); //
  Spawns an isolate from sample.dart file

  receivePort.listen((var message) {
    if(message is SendPort){
      sendPort = message;
      sendPort.send(receivePort.sendPort);
      sendPort.send("Hello");
      sendPort.send(["a", "list", "datatype"]);
    } else if(message is String) {
      print("Reply: $message");
    }
  });
}
```

Difference from Actor

Although, Dart isolates do not have shared state and use message-passing as the only means of communication between two isolates, the isolates differ from actors [section 2.2] in many ways. The most significant difference is the principle behind spawning

of actor and spawning of isolate. An actor is supposed to be a very lightweight and cheap to spawn but Dart isolates are resource heavy and slow to spawn. The implementation of actor found in other languages like Erlang [section 2.3] and Akka toolkit [section 2.5] can be considered much closer to Hewitt's actor model.

The number of actors that can be spawned per GigaByte of heap memory in Akka reaches up to 2.7 millions [Incb] whereas, an isolate in dart takes up around 5 to 7 MegaBytes⁶ of memory, the number of isolates per GigaByte of heap can only reaches up to few hundreds. Based on these observations, it would be appropriate to say that the current⁷ implementation of an isolate in Dart is — similar to a threads with properties like an actor.

Limitations of Isolates

Although Dart isolates follow the asynchronous message passing model which is suitable for distributed systems, it is not possible for an isolate in one Dart VM to send message to an isolate in another Dart VM. A message exchange between two isolates via SendPort/ReceivePort is possible only if the isolates are spawned locally in the same Dart virtual machine. Thus, a hindrance in making them distributed.

2.7 RabbitMQ - A Message Broker System

RabbitMQ is an open-source simple message broker software that implements Advanced Message Queuing Protocol(AMQP). It serves as an intermediary for message passing between applications or components. It give applications a common platform to send and receive messages, and keeps the messages safe until intended subscribers receive them [Sofe].

Messaging enables software applications to connect and scale. Applications can connect to each other, as components of a larger application, or to user devices and data. Messaging is asynchronous, decoupling applications by separating sending and receiving data. In RabbitMQ, messages are routed through exchanges before arriving at queues. RabbitMQ features several built-in exchange types for typical routing logic [Sofe].

RabbitMQ allows several servers of a local network to form a cluster. The cluster forms a single logical broker and queues are mirrored across several machines, which means the applications that uses it may connect to any of the servers that belong to the cluster [Sofe].

⁶Based on the memory consumption of the two isolates from Listing 2.3

⁷Dart version 1.7.2

RabbitMQ supports several protocols for enqueueing and dequeuing messages:

- AMQP (Several versions)

The Advanced Message Queuing Protocol (AMQP) was designed to provide reliability and interoperability. It provides messaging, including reliable queuing, topic-based publish-and-subscribe messaging, flexible routing, transactions, and security. AMQP exchanges route messages based on topic and headers[Pip]. Despite the fact that there are many different language implementations⁸ and examples for using AMQP in RabbitMQ, it is still not available for the Dart language.

- STOMP

STOMP [section 2.8] is a text-based messaging protocol emphasizing simplicity. More about STOMP is discussed in section 2.8.

RabbitMQ supports STOMP via a plugin – ‘rabbitmq_stomp’.

- MQTT

Message Queue Telemetry Transport was developed by IBM. It provides lightweight publish-and-subscribe messaging, targeted for resource devices with low resources and limited network bandwidth. Hence, the design principles and aims of MQTT are simpler and more focused than those of AMQP [Pip].

RabbitMQ supports MQTT 3.1 via a plugin.

- HTTP

HTTP is not a messaging protocol. Nevertheless, with the help of the listed technologies, that use HTTP as their substructure, RabbitMQ can transmit messages over HTTP [Soff].

Management Plugin

It supports a simple HTTP API to send and receive messages. This is primarily intended for diagnostic purposes but can be used for low volume messaging without reliable delivery.

Web-STOMP Plugin

The plugin supports STOMP messaging to the browser using WebSockets. For the older browsers that do not have support for WebSockets, fallback mechanisms provided by SockJS⁹ is used.

⁸Python, Java, Ruby, PHP, C#, Erlang etc. [Soff]

⁹SockJS is a JavaScript library that emulates WebSocket in browsers

JSON-RPC channel Plugin

This plugin support AMQP 0-9-1 messaging over JSON-RPC¹⁰. It is a synchronous protocol, thus the asynchronous delivery property of AMQP is emulated by polling.

2.7.1 Message Queues in RabbitMQ

In RabbitMQ, a queue is a mailbox name, where the messages are stored. It can buffer large quantity of messages bounded only by the available resources of the machine.

RabbitMQ receives messages from a client via one of protocols described in [section 2.7]. While sending a message, the client specifies the name of the queue where the message should be enqueued. Meanwhile, the subscribing client of the queue receives a message either as soon as the message is available in the queue, or when the client sends request for dequeue, or only after acknowledging previously dequeued message. This depends upon the parameters used during subscription to the queue.

If inflow of messages in a queue is faster than the outflow of the messages to its subscribers, both enqueueing as well as dequeuing gets slower. Since, messages are buffered into memory, as the quantity of messages increases, the consumption of the memory also increases. Besides, if there is a sudden increase in the inflow of messages, the incoming messages take more CPU time which results in the overall decrease of outflow of messages to consumers [Sofd].

2.7.2 RabbitMQ and prefetch-count

The prefetch-count Quality of Service (QoS) setting in RabbitMQ by default is set to unlimited. Which means, RabbitMQ empties the queue as fast as possible to the consumer. Setting the prefetch-count to unlimited might result in 'out of memory' or 'stack overflow' errors in the consumers. Again, setting the prefetch-count too low can hamper the performance of the whole application while setting it too high can cause the 'out of memory' exceptions. Hence, based on the requirement and design of the consumer application, appropriate prefetch-count should be determined.

The Figure 2.2 is the result of a benchmark¹¹ that shows how the throughput of a queue varies when prefetch count is changed for different number of consumers.

¹⁰JSON-RPC is JSON encoded Remote Procedure Call protocol

¹¹posted by Simon MacMullen on April 25th, 2012 at 2:47 pm




Figure 2.2: Chart¹² showing performance variation when prefetch count is changed in the consumer

2.8 STOMP

STOMP stands for Simple (or Streaming) Text Orientated Messaging Protocol. It is an alternative to other open messaging protocols such as AMQP. It provides an interoperable format for STOMP clients to communicate with a message broker system that supports STOMP. It provides interoperability among different languages, platforms and brokers [STO]. STOMP is designed to be a lightweight protocol that is easy to implement in both client and server.

2.8.1 Protocol Overview

STOMP is a frame based protocol. A frame consists of a command, a set of optional headers and an optional body. STOMP is text based but it allows the transmission of binary messages as well [STO]. A STOMP client can either be producer or consumer.

¹²Source: P. Software. *RabbitMQ Performance Measurements, part 2*. <http://www.rabbitmq.com/blog/2012/04/25/rabbitmq-performance-measurements-part-2/>. Last Accessed: 2014-11-20

2.8.2 STOMP Library for Dart

Dart's 'pub'¹³ is a public library sharing platform which allows user to share and reuse Dart codes.

Given that dart is a fairly new language, there is no AMQP client for RabbitMQ yet. As mentioned above in section 2.7 RabbitMQ also supports STOMP. An open source STOMP client¹⁴ in dart is available in Dart's 'pub' created by 'Potix corporation'. It can perform most of the basic operations with message broker system like connecting, creating queue, subscribing, enqueueing and dequeuing. Although it has those basic functionalities, it still has some limitations and incompleteness like lack of support of 'Heartbeat' and it only supports STOMP version 1.2 or above.

2.9 WebSocket

WebSocket protocol enables two-way communication between client and server over a single TCP connection. It uses origin-based security model, which is found in web browsers. It can be used for variety of web applications: games, stock tickers, multiuser applications, user interfaces exposing server-side services in real time, etc [FM11].

Since, HTTP was not initially designed for bidirectional communication, the WebSocket Protocol is designed to displace other existing bidirectional communication technologies that are based on HTTP [FM11].

WebSocket uses two URI schemes: "ws://" for normal WebSocket connection and "wss://" for secured WebSocket connection [FM11].

2.9.1 Security

The WebSocket Protocol uses the origin model used by web browsers to restrict which web pages can contact a WebSocket server when the WebSocket Protocol is used from a web page [FM11].

A WebSocket server reads the handshake sent by client to establish a connection. Thus, an attempt to connection to WebSocket from other protocols cannot succeed if it is not sent by a WebSocket client [FM11].

2.9.2 Establishing a Connection

When establishing a WebSocket connection, the HTTP server receives a regular GET request with an offer to upgrade to WebSocket. The server responds to the request to

¹³<http://pub.dartlang.org>

¹⁴<https://github.com/rikulo/stomp>

complete the handshake and establish the connection. Then the communication takes place in full-duplex mode [FM11].

3 System Design

3.1 Core Design Decisions

The DDE Framework is designed to be distributed in nature with the concept of Actor Programming [section 2.2]. It follows the standard actor programming concept and provides inherently distributed nature to the applications built on top of it. The framework itself is built using the concept of ‘message-passing’ [section 2.1] to alleviate any possibility of concurrency issues, thus making the applications thread-safe.

- The framework does not guarantee the delivery of a message.
- All the messages sent by the framework are based on ‘fire and forget’ concept.
- A message is delivered at most once.
- For consistency and persistence, a message is always routed through Message Queuing System [subsection 3.2.3], even though the target isolate may belong to same isolate system in the same logical or physical node.
- Dequeueing a message (from MQS) by an isolate is based on pull mechanism, not push mechanism.
- Exceptions thrown at child isolates are handled by ‘spawner’ of that isolate. Hence, implementing the idea of supervision and “let it crash” ideology [subsection 2.3.1].

3.2 The Framework




Figure 3.1: Architecture of the framework

The framework comprises of an Isolate System, a Registry, a Message Queuing System, a Message Broker System and an Activator. The Figure 3.1 depicts the general overview of different components and sub-components of the framework.

3.2.1 Isolate System

An Isolate System is analogous to an actor system subsection 2.5.2. Just as an actor system consists of a group of actors working together, an isolate system is composed of a group of 'Worker' isolates. It consists of different hierarchies which forms an organizational-like structure. The top level isolate is the Isolate System itself and the bottom most are 'Worker' isolates.

A 'Bootstrapper' in a physical node can start up several Isolate Systems. Nevertheless, a logical Isolate System is not limited to a single physical node. The 'Worker' isolates spawned by an isolate system can be distributed across several remote systems.

Each isolate system has a unique id, which is a UUID. It is generated when the isolate system is bootstrapped. For bootstrapping, an isolate system needs the WebSocket address of Message Queuing System, and a 'name' for itself. The name is simply an alias, and should not be confused with the unique id as another isolate system with the same name can exist in other nodes but the unique id is exclusive for a particular instance of an isolate system.

The bootstrapping of an isolate system includes: generating a new id which is unique for itself, opening up a 'ReceivePort', and connecting to a 'Message Queuing System'. After opening up a 'ReceivePort', the isolate system starts listening on that port for messages so that it can receive incoming messages from 'Controller'. While connecting to the Message Queuing System, if a connection could not be established, it simply keeps on retrying at certain interval. Furthermore, should the connection be lost at any time after being connected with the MQS, the Isolate System automatically keeps on trying to re-establish the connection at regular intervals. Since, the connection to MQS takes place asynchronously, the isolate system moves forward and spawns a controller, regardless of the establishment of connection to the MQS.

Adding Worker Isolates to The Isolate System

As an isolate system is a top level isolate, it spawns controller. The controller spawns one or several routers and each router spawns a worker isolates. Figure 3.2.1 shows the message flow sequence in different components while starting up an isolate system and deploying a worker isolate in it.




Figure 3.2: Adding a Worker isolate to an isolate system

When an isolate system is first initialized, it is an empty system without any Workers running in it. The Workers can be added with appropriate load balancer once the isolate system has been initialized. The 'addIsolate' method starts up worker isolates into the isolate system. It requires following arguments:

name – A name for pool of isolates. A deployed isolate has its own name but overall the name is concatenated with the name of isolate system to denote the hierarchy. For instance, an isolate with name 'account' becomes 'bank/account' where 'bank' is the name of the isolate system.

sourceUri – The location of the source code from which the isolate shall be spawned. The path can either be absolute path to the local file system or the full http or https URI.

workersPaths – List of destinations where each of the isolates should be spawned. To spawn locally, 'localhost' should be used, whereas to spawn in remote node, WebSocket path like: "ws://192.168.2.9:42042/activator" should be used. The number copy of isolates that should be spawned is determined by the length of

this list. If multiple copies of isolates should be spawned in a node, the location can be repeated. For instance [“localhost”, “localhost”] results in spawning of two identical isolates in local machine which is load balanced by the type of specified router.

routerType – The type of load balancing technique that one would like to use to effectively distribute incoming messages. By default, the framework, provides three types of routers: Round-Robin, Random and Broadcast. If the developer wants to use his own customized load balancer instead of using the options provided by the framework, the absolute path to the location of source code, which can also be a remote URI, of the custom router implementation can be provided.

hotDeployment – This argument is optional and is set to ‘true’ by default. Setting it to ‘true’ enables continuous monitoring of the source code. If any change in source code is identified, the instances of isolates, spawned by this ‘addIsolate’ function, in current isolate system will be restarted, without the need of redeploying the system.

args – Custom additional arguments to be passed into each instance of spawned isolate. This argument is also optional and can be safely ignored.

Message Handling in The Top Level Isolate System

Typically, a message in an isolate system may arrive from three sources: Message Queuing System via WebSocket, Controller via ReceivePort or Bootstrapper via direct method invocation. As Dart is a single threaded programming language [subsection 2.6.5], only one message is handled at a time by the top level isolate.

A Message Queuing System sends the data over WebSocket in JSON string format, which should be deserialized to Map data type before further processing. As the received message contains the queue name from which it is dequeued, the name of the queue is then parsed and transformed to name and address of the corresponding isolate. The message is then forwarded to the Controller that this instance of the isolate system has spawned.

The messages arriving from a controller are either dequeue requests or messages for enqueueing (that should be sent to the Message Queuing System). The dequeue requests are sent from the isolates that have completed certain task and are ready to accept another message. For the dequeue requests, the sender of the message is identified, which is used to figure out the corresponding name of the queue name. Then the pull request to dequeue from that queue is forwarded to MQS via open WebSocket port.

For the messages that are supposed to be delivered to another isolate, the name of the target isolate is used to figure out the name of the queue and then sent to the MQS for enqueueing.

The 'Bootstrapper' of a node that creates an isolate system can send messages to isolate system by directly invoking the functions provided by the isolate system. The bootstrapper can request the information about the isolates this instance of isolate system is running. For which, the isolate system delegates the message to its controller and waits asynchronously for the response from the controller. The request, to fetch a list of running worker isolates, is triggered when a user sends the request to view details of an isolate system via a web interface or via RESTful web services provided by the 'Registry' [subsection 3.2.2].

Another type of message is the 'KILL' message, which is used to terminate a worker isolate. It is also forwarded to the controller as the isolate system does not directly manage the running worker isolates. Thus, the message is forwarded to the controller which is next in the hierarchy. In contrast, when the shutdown command for the isolate system is triggered via web or REST interface, the isolate system closes all the open ports including WebSocket ports and ReceivePorts, and then wait for the 'Garbage Collector' to clean up the memory reserved by it.

Controller

Every isolate system has a single controller, which is spawned by the top level isolate of the isolate system. A controller stays idle until it receives a message to create an isolate. Basically, a controller spawns and manages all the routers of an isolate system. Additionally, a controller takes care of the 'hot deployment' feature for which it spawns a 'FileMonitor' for each router if the feature is enabled. When a RESTART message is received from a FileMonitor, the controller sends a RESTART_ALL message to the designated router, which restarts all the Worker isolates the router has spawned.

A controller is also responsible for replying to the query of list of isolates an isolate system is running. It achieves this by keeping a detailed record of each Router and number of Worker isolates each Router is handling, which is updated as soon as an isolate is killed or a new isolate is added.

As a Controller is the 'spawner' of Routers and the 'spawnee' of the top level isolate, it forwards the messages as well as dequeue requests coming from Routers to the top level isolate of an isolate system.

Router

A router is spawned by a controller. The router creates and is responsible for a group of identical Workers isolates. Since an isolate is single threaded, creation of multiple instances of an isolate is desirable for concurrency. When a message arrives in a router from a controller, the router, based on its defined routing policy, delegates the message to one of the worker isolates. The routing policy can be chosen at the time of deployment of a worker isolate.

A router uses a routing policy to distribute message among the group of isolates it is handling. The default routing policies that are available in the framework are listed in the Table 3.1

Table 3.1: List of routing techniques provided by the framework

Router	Description
Round Robin	Messages are passed in round-robin fashion to its Worker isolates.
Random	Randomly picks one of its Worker isolates and sends the message to that Worker isolate.
Broadcast	Replicates and sends message to all of its Worker isolates.

In addition to the available routing policies of the framework, it is also possible to add a new Routing technique by simply extending the 'Router' class which requires 'selectWorker' function to be implemented. The overridden 'selectWorker' function may either return a list of Workers or a single Worker. The ability to implement a custom router opens up possibilities for numerous load balancing techniques. For instance, a simple multicasting router that replicates a message only to the Workers that are spawned locally can be implemented by selecting such Workers using their deployment paths and returning them as a 'List'.

As the router manages the Worker isolates it has spawned, it is responsible for effectively terminating and restarting the Worker isolates. It also buffers the messages that might arrive while the workers are not ready to accept the messages yet; usually, during the creation of Worker isolates and while restarting them.

If a router does not receive any message from a Worker for a certain amount of time, the router sends a PING message to check if the Worker isolate is alive and ready to accept more messages. If the Worker isolate responds with a PONG message, the router sends a request to fetch messages to controller. This mechanism is present in the framework to prevent the 'starvation' for a Worker isolate in case the dequeue message, that might have been sent earlier, could not reach the Message Queuing System because

of a network issue or unavailability of MQS.

Worker

The 'Worker' of the framework is an abstract class, which should be inherited by the isolate that the programmer creates. The 'Worker' first unwraps the messages that arrives from the router and retrieves headers from it. 'Sender' and 'replyTo' headers of the message are collected before forwarding message to the child class, that extends this abstract class. By unwrapping the messages that are encapsulated by various headers, the abstract Worker class makes sure that the messages are delivered to the target implementation of the Worker isolate immutated and in intended form.

To extend the 'Worker' isolate, one must implement 'onReceive' function which handles incoming messages and carry out the business logic tasks. However, if a task is too complex, the Worker isolate can divide the tasks into subtasks and spawn temporary isolates to carry out those subtasks concurrently. The temporary isolates can be terminated once the subtask has been carried out.

The 'send', 'reply' and 'ask' functions are provided by this abstract Worker class to send a message to another worker isolate. These functions automatically add the information of sender and receiver in the header of the message that is sent out.

Sending a Message To send a message from one Worker isolate to another, the framework provides 'send' function. It takes 'message' and 'address' of the target Worker isolate as its argument. The reply path can also be optionally set, so that the replied message from target isolate is sent to a different worker isolate for further processing. The named parameter¹ 'replyTo' can be used to set the address of intended recipient of the replied message; eg:

```
|| send("A simple text message", "demosystem/printer");  
|| send("Another message", "demosystem/jsonConverter", replyTo: "demosystem/printer");
```

Asking For a Reply Sometimes a Worker isolate might need a reply from another isolate for further processing or before replying to the sender of the message. In such case, the worker isolate can specifically ask the target isolate to reply to this particular instance of worker isolate.

For instance, a sample use case can be, a worker isolate maintaining a connection with a browser via HTTP. In this case, as the port cannot be serialized and passed to other isolates through messages, another instance of similar isolate will not be able to respond to the request made in that connection.

¹Dart's named paramter is an optional argument in the function

Similar to the 'send' function, the 'ask' function takes 'message' and 'address' of the target worker isolate as its argument; eg:

```
|| ask("current time", "demosystem/timeKeeper");
```

Replying to a Message To reply to a message, the framework provides the 'reply' function. It expects a single argument — 'message', because the response is sent to the worker isolate specified by the sender. The 'reply' can be used in response to any of 'send' or 'ask' messages; eg:

```
|| reply("Current time is: $time");
```

Proxy

A 'Proxy' is a special type of a Worker isolate. When a Worker isolate is supposed to be spawned in a remote node, the router instead spawns a Proxy isolate in local node. Once the Proxy isolate is created, it connects to the 'Isolate Deployer' of the remote node where the Worker isolate is intended to be spawned. After establishing connection over a WebSocket with Isolate Deployer of the remote node, the proxy isolate forwards the request to spawn the worker isolate to the Isolate Deployer. After successful spawning of isolate in the remote node, the proxy isolate simply forward the messages that are sent to it by the spawner router. Each proxy worker maintains a separate WebSocket connection with an 'Isolate Deployer'.




Figure 3.3: A Proxy Worker

FileMonitor

The controller spawns a 'FileMonitor' only if the 'Hot Deployment' flag for a worker isolate is set while deploying. The spawned 'FileMonitor' monitors the md5 checksum of the source code from which the Worker isolate is spawned. If a change in the source file is detected, it simply sends a RESTART command to the controller, which eventually forwards it to the target router. The router then restarts all of its worker isolates.

3.2.2 The Registry

The Isolate Registry is a central node where other nodes, which are running Bootstrapper, connect and register themselves. The registry simply keeps the record of the connected nodes, assigns a unique id to each and queries them about the running isolate systems when required. The registry provides RESTful API and a web interface ² through which one can have an overview of the full system and manage the deployments of the isolate systems as well as individual isolates.

The basic tasks that a registry carries out are:

- Bootstraps an isolate system, during runtime, in local or remote node
- Provides a way to deploy, update or remove an isolate system
- Returns information about the deployed isolates by querying the individual isolate system of a node.

RESTful API of Registry

The registry provides a REST API to perform the operations on the connected nodes. One can send a 'GET' request to the registry to fetch the list of the nodes that are connected to the registry. Using the 'id' of a node from the replied list, one can deploy an isolate system or add an isolate to already deployed system by sending appropriate 'POST' request.

The REST endpoints exposed by the registry are listed in table Table 3.2.

²the web interface can be opted out as it has to be started separately

Table 3.2: List of endpoints exposed by the registry

Method	Endpoint
GET	/registry/system/list
GET	/registry/system/{bootstrapperId}
POST	/registry/deploy
POST	/registry/system/shutdown

A sample GET and POST query to deploy an isolate system

GET request to fetch a list of connected systems

Request: 'GET http://54.77.239.254:8000/registry/system/list'

Response Status Code: 200 OK

Response Body:

```

1  [
2    {
3      "bootstrapperId": "266393094",
4      "ip": "54.77.239.244",
5      "port": "50189"
6    },
7    {
8      "bootstrapperId": "12133208",
9      "ip": "54.77.239.243",
10     "port": "50192"
11   }
12 ]

```

POST request to deploy an isolate system

Request: 'POST http://54.77.239.254:8000/registry/deploy'

Request Body:

```

1  {
2    "bootstrapperId" : "266393094",
3    "action": "action.addIsolate",
4    "messageQueuingSystemServer": "ws://54.77.239.200:42043/mqs",
5    "isolateSystemName" : "sampleSystem",
6    "isolateName" : "consumer",
7    "uri" : "http://54.77.239.221/sampleSystem/bin/Consumer.dart",

```

```
8  "workersPaths" : ["localhost",
9    "ws://54.77.239.243:42042/activator"],
10 "routerType" : "random",
11 "hotDeployment" : true
12 }
```

Response Status Code: 200 OK

A sample GET query to fetch details of an isolate system

GET request to get details of an isolate system

Request: 'GET http://54.77.239.254:8000/registry/system/266393094'

Response Status Code: 200 OK

Response Body:

```
1  {
2    "sampleSystem": [
3      {
4        "id": "sampleSystem/consumer",
5        "workerUri":
6          "http://54.77.239.221/sampleSystem/bin/Consumer.dart",
7        "workersCount": 2,
8        "workersPaths": [
9          "localhost",
10         ws://54.77.239.243:42042/activator"
11       ],
12        "routerType": "random",
13        "hotDeployment": true
14      }
15    ]
16  }
```

A sample POST query to terminate a Worker isolate

POST request to terminate an isolate of an isolate system

Request: 'POST http://54.77.239.254:8000/registry/system/shutdown'

Request Body:

```
1  {
2    "bootstrapperId" : "266393094",
```

```
3  "isolateSystemName" : "sampleSystem",  
4  "isolateName" : "consumer"  
5  }
```

Response Status Code: 200 OK

An example of terminating an Isolate System

POST request to terminate an isolate of an isolate system

Request: 'POST http://54.77.239.254:8000/registry/system/shutdown'

Request Body:

```
1  {  
2    "bootstrapperId" : "266393094",  
3    "isolateSystemName" : "sampleSystem"  
4  }
```

Response Status Code: 200 OK

The registry generates all the information about isolate and isolate systems “on the fly”. Thus, it does not need to persist any data.

The Web Interface for the Registry

The deployment of isolates can also be managed by using a web interface provided by the registry. The Web Interface should be started up separately in a different port. The Web Interface, internally communicates with the registry via the REST API [3.2.2] which is exposed by the Registry.

3.2.3 Message Queuing System (MQS)

Since, the basis of this system is message passing, the Message Queuing System is an important component of this framework. The MQS is a top level isolate that fetches messages from message broker system and dispatches to respective isolate of the isolate system. Whenever a new isolate system starts up, it opens up a new WebSocket connection with the MQS. The messages are exchanged between the isolate system and the MQS through a WebSocket connection. The MQS keeps track of the unique-id of an isolate system so that it can identify the origin of the message.

If a message is supposed to be enqueued, the MQS ignores the unique-id and simply forwards the message to the ‘Enqueuer’ isolate. Whereas, if the message is a dequeue request, the MQS forwards the the message to a ‘Dequeuer’ isolate along with the

unique-id of the isolate system. The unique-id is used to identify the WebSocket port through which the request arrived. Thus, making sure that the dequeued message is sent to the correct requester. This is required if a cluster, of identical isolate systems, is running on different nodes.

The MQS should be started up separately along with few command line arguments to connect to message broker system. The required command line arguments are: ip address, port, username and password to connect to Message Broker System. The 'prefetchCount' is an optional argument which defaults to 1, if not provided explicitly. A 'prefetchCount' is a Quality of Service header for Message Broker System which allows a subscriber of a queue to hold the defined quantity of unacknowledged messages.

Since, in Dart³ the passing of sockets to isolates is not yet possible, the main isolate has to pipe all the input/output data. In this case the MQS is the top level isolate which has to handle all incoming and outgoing messages.

Enqueuer

An enqueuer is a separate isolate. A Message Queuing System has only one enqueuer, which basically receives messages from the MQS and sends messages to a message broker system – RabbitMQ [2.7] via STOMP [2.8] protocol.

Dequeuer

As opposed to Enqueuer, a Message Queuing System maintains each dequeuer for each topic. The topic corresponds to each router running in the isolate system. Whenever a message arrives from a new isolate, the MQS spawns a new dequeuer isolate. The dequeuer then subscribes to a new message queue in the message broker system via STOMP [2.8] protocol. If the queue does not exist in the message broker system, the message broker system automatically creates the queue.

If a dequeuer is idle for too long, i.e. if the Dequeuer isolate has not received any dequeue requests for certain interval⁴, then the MQS terminates the dequeuer isolate for that particular queue. Nevertheless, as soon as the MQS receives a dequeue request, it spawns a new Dequeuer, if one does not exist yet.

The dequeuer subscribes messages from Message Broker System with such options that the new messages do not arrive to the subscriber unless previously dequeued messages have been acknowledged. This throttles the flow of messages from message broker system and keeps itself and the isolates from being overwhelmed by a large number of messages, which might induce 'out of memory' issues.

³Dart version 1.7.2

⁴by default the timeout is 10 seconds

Messages in dequeuer keep on arriving as long as there are messages in the queue and the messages are being acknowledged. The messages that are in the buffer of dequeuer stay in unacknowledged state unless they are flushed and sent out to the requesting isolate of an isolate system. As soon as a message is acknowledged the dequeuer receives another message from the message broker.

Multiple Instances of MQS

It is possible for a system to have multiple Message Queuing Systems for scaling up the system. If there are multiple identical isolate systems connected to different instances of MQS, each MQS will have a dequeuer which subscribes to the same queue. Nevertheless, a message is dispatched by the message broker system to only one of the dequeuers, which is distributed in round robin fashion. Thus, messages are fairly distributed among the subscribers.

3.2.4 Activator

An activator simply starts up two isolates: a 'System Bootstrapper' and an 'Isolate Deployer'. Every node that is supposed to be running an Isolate System or become a part of isolate system by running isolates must be running an Activator. The activator requires a WebSocket address of the Registry as a command line argument. Nevertheless, it is also possible to start up System Bootstrapper and Isolate Deployer separately.

System Bootstrapper

The System Bootstrapper registers itself to the 'Registry' via a WebSocket connection as soon as it is started. The activator forwards the path of the WebSocket to the system bootstrapper. But, if a System Bootstrapper is started separately then the path of the Registry should be passed as a command line argument.

Isolate Deployer

An Isolate Deployer starts up a Worker isolate in a node. The isolate is spawned without a local isolate system and as a part of an isolate system running in another node. This functionality expands the isolate system beyond a physical system. An isolate system can deploy number of instances of an isolate in several different nodes.

An isolate deployer running in a remote machine is able to handle requests from multiple 'Proxies' from several isolate systems. Each proxy opens up a separate WebSocket channel with the isolate deployer.

3.3 Some Key Features

3.3.1 Hot Deployment of Isolates and Isolate Systems

It is possible for the source code, of an isolate, to reside in a remote repository and fetched by the controller of a node when required. For instance: isolate source code can reside in a git repository hosted in GitHub. So that as soon as new code is committed in the repository, it gets immediately picked up by the application and the change gets reflected without restarting the application.

After a node is bootstrapped, changes like: addition, update or removal of isolates in an isolate system can take place. In such case, the isolates can be killed and redeployed when it has finished processing tasks and is sitting idle. A dedicated isolate 'FileMonitor' [section 3.2.1] monitors changes in the code repository. When a change is detected, the 'FileMonitor' isolate sends a RESTART message along with the target router to notify the controller. The controller takes care of pushing the message to relevant router, and the router takes care of terminating and re-spawning the worker isolates.

This hot deployment capability improves the availability of an application. Whenever there is any change in a component of an application, the whole application does not need to be re-deployed, instead, only a set of isolates that should be updated is restarted at runtime. This increases overall up-time of the application and keeps other components working even in the time of modifications.

3.3.2 Migration of Isolates and Isolate Systems

Relocation of Worker isolates or an isolate system during runtime i.e. killing a set of Worker isolates or an isolate system at one node and bringing up same set of Worker isolates in another node is the migration of isolates or isolate system. The concept of hot deployment and migration brings enormous possibilities in a distributed system. Some of them are:

- Migration of actors/isolates allows an application to scale in an easy way. With this capability, it is also possible for an application to be brought up the most frequently used isolates near to the server where it is accessed the most.
- Related and dependent isolates can be migrated to the same server, if it is evident that it improves performance of the entire system.
- In case of hardware failure on a system which is running a certain set of isolates, migration of worker isolates during runtime can make the application survive the hardware failure.

3.3.3 Remote Isolates

The current isolate implementation in Dart⁵ cannot communicate with other isolates over a network. The Worker isolates in this framework have the ability to communicate with the isolates that may be running in a remote node. So, there can be isolates running in any node. The communication underneath is taken care of by the framework so the developer using this framework does not have to worry if an isolate is remotely spawned or locally spawned.

Two isolates, although, running in two different virtual machines, can still belong to a same logical isolate system.

3.4 Typical Message Flow in the System

The framework is based on ‘fire-and-forget’ principle of message sending. The Figure 3.4 shows a simple message flow while enqueueing a message and the Figure 3.5 shows the process of dequeuing a message. A message is serialized to JSON string before sending via a SendPort of an isolate and deserialized after receiving from ReceivePort.




Figure 3.4: Enqueueing a message

⁵Dart version 1.7.2




Figure 3.5: Dequeuing a message

3.4.1 Sample message formats

The message sent through different components while enqueueing

Original Message:

```
"Test"
```

Worker:

```
{senderType: senderType.worker, id:
  sampleSystem/producer/88f52440-5060-11e4-f396-97cebb949945,
  action: action.send, payload: {sender: sampleSystem/producer,
  to: sampleSystem/consumer, message: Test, replyTo: null}}
```

Router:

```
{senderType: senderType.router, id: sampleSystem/producer, action:
  action.send, payload: {sender: sampleSystem/producer, to:
  sampleSystem/consumer, message: Test, replyTo: null\}\}}
```

Controller:

```
{senderType: senderType.controller, id: sampleSystem/producer,
  action: action.send, payload: {sender: sampleSystem/producer,
  to: sampleSystem/consumer, message: Test, replyTo: null\}\}}
```

Top level isolate:

```
{targetQueue: sampleSystem.consumer, action: action.enqueue,
  payload: {sender: sampleSystem/producer, message: Test, replyTo:
  null}}
```

Message Queuing System:

```
{topic: sampleSystem.consumer, action: action.enqueue, payload: {
  sender: sampleSystem/producer, message: Test, replyTo: null}}
```

Enqueuer:

```
{sender: sampleSystem/producer, message: Test, replyTo: null}
```

Sample format of a message sent at different components while dequeuing

Dequeuer:

```
{"sender": "mysystem/producer", "message": "Test", "replyTo": null}
```

Message Queuing System:

```
{senderType: senderType.dequeuer, topic: mysystem.consumer, payload:
  {sender: mysystem/producer, message: Test, replyTo: null}}
```

Top level isolate of isolate system:

```
{senderType: senderType.isolate_system, id: mysystem, action:
  action.none, payload: {to: mysystem/consumer, message: {sender:
  mysystem/producer, message: Test, replyTo: null}}}}
```

Controller:

```
{senderType: senderType.controller, id: controller, action:
  action.none, payload: {to: mysystem/consumer, message: {sender:
    mysystem/producer, message: Test, replyTo: null}}}}
```

Router:

```
{senderType: senderType.router, id: mysystem/consumer, action:
  action.none, payload: {sender: mysystem/producer, message: Test,
    replyTo: null}}
```

Worker:

```
"Test"
```

3.4.2 Some Implementation Overview

Some insight about the implementation of selected functions of the framework:

How 'send' works?

When a Worker Isolate sends a message using 'send' function of the Worker class, the message is encapsulated with further information about the sender and the receiver are added to the message. The encapsulated message is forwarded to the spawner isolate, which in this case, is the router. The router again forwards it to the controller which again forwards to the top level isolate. Then the top level isolate adds another level of encapsulation and headers to the message so that the Message Queuing System knows the destination queue.

If the Worker isolate is expecting to consume another message after sending a message, it should send a PULL Request for another message, which can be performed by invoking the 'done' function.

How 'ask' works?

Ask function has certain subtle differences from the 'send' function. The 'ask' function should be used when the sender of the message expects another message in reply. The abstract Worker class adds the full path of the isolate along with the unique-id of the isolate when an 'ask' message is constructed. This is to make sure that the response from the target isolate reaches this particular instance of the isolate. When the router receives a message with complete address (name along with its unique id) of a Worker isolate, it routes the message to the specific worker regardless of its routing algorithm. If the worker isolate with given unique

id is not found in the list of worker isolates the router is maintaining, then the message is simply discarded by the router. This case is possible when the isolates have been restarted or for some reason the isolate was killed.

How 'reply' works?

The 'reply' function is simply a convenience for the implementer. The 'reply' function simply invokes 'send' message with the sender's address as the target isolate. If the message contains 'replyTo' then the message will be replied to the address contained in 'replyTo' instead of the original sender. The 'reply' function can be used to reply message in both – 'send' and 'ask' cases.

How 'KILL' works?

This is a special control message sent to the isolates as well as isolate system to shutdown themselves. If a KILL message is sent to a Worker isolate, the message is enqueued to the end of the Worker isolate as any other message. Messages arriving after a KILL message are not sent to the worker isolate and are buffered until the isolate is restarted. The messages are delegated to the worker isolate only after its spawning is complete.

When a Worker isolate finishes processing queued messages and encounters the KILL message, the isolate closes its ReceivePort⁶ and stays idle. After sometime it gets 'Garbage Collected' and is cleaned up. But, sometimes the garbage collector cannot clean up the isolate and a memory leak occurs. Thus, as a workaround, a custom Exception is thrown deliberately by the isolate to terminate itself. The exception is thrown only after it closes all ports. This workaround forcefully terminates the isolate and releases the memory occupied by the isolate.

The abstract Worker class provides *beforeKill()* method, which can be overridden to perform custom operations before terminating an isolate.

How RESTART works?

Restarting an isolate is basically a combined process of killing an isolate and spawning it up again. However, during the restart, after issuing the KILL message, the messages may keep coming from the controller to the router. These messages are buffered in the router itself. The buffered messages are flushed and sent out once the Worker isolates are spawned. For instance, when the 'Hot Deployment' subsection 3.3.1 feature is enabled, if the source code of the isolate is modified and saved, each of the isolates that the router has spawned gets restarted. During which the messages that arrive after RESTART message are buffered in the router.

⁶A Worker Isolate receives message from Router via a ReceivePort

How shutting down an isolate system works?

An isolate system that is running in a node can be shutdown via the Web Interface or via POST request to the registry. When a request to shutdown an isolate system is sent, the isolate system closes all the ports including the isolate ports as well as the WebSocket connection with Message Queuing System. After that the forceful shutdown is carried out by throwing out a custom Exception. This is a work-around to free up the memory consumed after it is shutdown, because the feature to immediately terminate an isolate is yet to be implemented in dart.

3.4.3 Clustering

Clustering can be achieved in the framework in several levels.

- By deploying worker isolates in several remote nodes. i.e. taking advantage of the concept of 'Remote Isolates' [subsection 3.3.3].
- By deploying replicas of an isolate system in different nodes. An isolate system with same name can exist in another node even though they connect to the same MQS.
- The Message Queuing System itself, can also be replicated where replicas of isolate system may connect to different instance MQS.
- Since, several instances of RabbitMQ [section 2.7] can form a logical group, sharing common configuration, properties, users, queues etc., a cluster of message broker system can be formed. Which allows, Message Queuing Systems to connect to the different member of a cluster.

<TODO: a simple diagram>

3.5 Dart Libraries Used in Construction

Table 3.3: List of libraries directly used by the framework

Library	URL
path	https://pub.dartlang.org/packages/path
uuid	https://pub.dartlang.org/packages/uuid
crypto	https://pub.dartlang.org/packages/crypto
stomp	https://pub.dartlang.org/packages/stomp

3.6 A Sample Implementation of Worker Using The Framework

Listing 3.1 is an example implementation of Worker isolate for the framework. The *main()* is the entry point of the isolate, which is invoked when this isolate is spawned by the router. The class *Consumer* overrides the function *onReceive()*, which is called for each incoming message.

The example shown here prints a message, if the 'action' set in the *message* variable is "print". If the 'action' set in the *message* variable is "send_back" then the Worker isolate sends the message back to the sender of the message.

Listing 3.1: A sample Worker isolate that can be deployed in the framework

```
import 'dart:isolate';
import 'package:isolatesystem/worker/Worker.dart';

main(List<String> args, SendPort sendPort) {
  Consumer printerIsolate = new Consumer(args, sendPort);
}

class Consumer extends Worker {
  Consumer(List<String> args, SendPort sendPort) : super(args, sendPort);

  @override
  onReceive(message) {
    switch(message['action']) {
      case "print":
        print("message['content']");
        break;
      case "send_back":
        reply(message['content']);
        break;
    }
    done();
  }
}
```

4 Results

4.1 Applications based on the DDE framework

To test the throughput of messages, scalability of the framework and latency of message delivery two sample applications were created based on the DDE framework. The sample applications subsection 4.1.1 and subsection 4.1.2 were used for testing and evaluation of the framework.

4.1.1 Producer/Consumer Program

The producer-consumer program is a simple application. It has two worker isolates, a producer and a consumer. The 'producer' shown in Listing 4.1 produces messages continuously until it is terminated. The messages are targeted for the 'consumer' [Listing 4.2] which prints the incoming messages to standard output.

Listing 4.1: Basic version of Producer Worker of Producer-Consumer application

```
// producer.dart
import 'dart:io';
import 'dart:isolate';
import 'dart:async';
import 'dart:math' as Math;

import 'package:isolatesystem/worker/Worker.dart';

main(List<String> args, SendPort sendPort) {
  Producer producer = new Producer(args, sendPort);
}

class Producer extends Worker {
  static const String consumerAddress = "mysystem/consumer";
  static const String Message64Bytes = "012345670123456701234567";

  StringBuffer data = new StringBuffer();
  String description = "";
```

```
Producer(List<String> args, SendPort sendPort) : super(args, sendPort) {
    description += "${args}";
    sendMsgWithDelay();
}

@override
onReceive(message) {}

sendMsgWithDelay() {
    while(true) {
        int timestamp = new DateTime.now().millisecondsSinceEpoch;
        Map message = {'createdAt': timestamp, 'message': Message64Bytes};
        int delay = 20 + new Math.Random().nextInt(480);
        sleep(new Duration(microseconds:delay));
        send(message, consumerAddress);
    }
}
}
```

Listing 4.2: Basic version of Consumer Worker of Producer-Consumer application

```
//consumer.dart
import 'dart:isolate';
import 'package:isolatesystem/worker/Worker.dart';

main(List<String> args, SendPort sendPort) {
    Consumer printerIsolate = new Consumer(args, sendPort);
}

class Consumer extends Worker {
    Consumer(List<String> args, SendPort sendPort) : super(args, sendPort);

    @override
    onReceive(message) {
        outText(message);
    }

    outText(var message) {
        print(message);
        done();
    }
}
```


4.1.2 Requester/Supplier Program

The request-reply program is a simple application. It has two worker isolates, a requester and a supplier. The 'requester' shown in Listing 4.3 produces a request message as soon as it is spawned. The request message is targeted for the 'supplier' [Listing 4.2] which does a pattern matching on the request and replies an appropriate message to the original requester.

Listing 4.3: Requester Worker of Requester-Supplier application

```
import 'dart:isolate';
import 'dart:math';

import 'package:isolatesystem/worker/Worker.dart';

import 'AvailableSupplies.dart';

main(List<String> args, SendPort sendPort) {
  new Requester(args, sendPort);
}

class Requester extends Worker {
  static const String SupplierAddress = "demosystem/supplier";
  DateTime startTime;

  Requester(List<String> args, SendPort sendPort):super(args, sendPort) {
    startTime = new DateTime.now();
    _sendRequest();
  }

  @override
  onReceive(message) {
    int receivedAt = new DateTime.now().millisecondsSinceEpoch;
    int requestedAt = message['requestedAt'];
    int repliedAt = message['repliedAt'];

    int rtt = receivedAt - requestedAt;
    print("Round trip time: $rtt");
    _sendRequest();
  }

  _sendRequest() {
    String requestMessage = "";

    int randomNumber = new Random().nextInt(3);
    if(randomNumber == 1) {
      requestMessage = AvailableSupplies.RANDOM_FRUIT;
```

```
    } else if (randomNumber == 2) {
      requestMessage = AvailableSupplies.RANDOM_NUMBER;
    } else {
      requestMessage = AvailableSupplies.RANDOM_NAME;
    }

    Map message = {'requestedAt':new DateTime.now().millisecondsSinceEpoch,
      'requestMessage':requestMessage};
    ask(message, SupplierAddress);
  }
}
```

Listing 4.4: Requester Worker of Requester-Supplier application

```
import 'dart:isolate';
import 'dart:math';

import 'package:isolatesystem/worker/Worker.dart';
import 'AvailableSupplies.dart';

main(List<String> args, SendPort sendPort) {
  new Supplier(args, sendPort);
}

class Supplier extends Worker {
  Supplier(List<String> args, SendPort sendPort):super(args, sendPort);

  @override
  onReceive(message) {
    String responseMessage = "";
    switch(message['requestMessage']) {
      case AvailableSupplies.RANDOM_FRUIT:
        responseMessage = _randomFruit();
        break;
      case AvailableSupplies.RANDOM_NUMBER:
        responseMessage = _randomNumber();
        break;
      case AvailableSupplies.RANDOM_NAME:
        responseMessage = _randomName();
        break;
    }
    Map response = {
      'requestedAt': message['requestedAt'],
      'repliedAt': new DateTime.now().millisecondsSinceEpoch,
      'requestMessage': message['requestMessage'],
      'responseMessage': responseMessage};
  }
```

```
        reply(response);
        done();
    }

    String _randomFruit() {
        int number = new Random().nextInt(5);
        return ["APPLE", "ORANGE", "KIWI", "PEAR", "GRAPES"][number];
    }

    String _randomNumber() {
        return new Random().nextInt(99999).toString();
    }

    String _randomName() {
        int number = new Random().nextInt(5);
        return ["Lorna", "Ambrose", "Domingo", "Kirsten", "Zachery"][number];
    }
}
```

Listing 4.5: Supporting class that contains list of constants for pattern matching

```
class AvailableSupplies {
    static const String RANDOM_FRUIT = "supply.fruit";
    static const String RANDOM_NUMBER = "supply.number";
    static const String RANDOM_NAME = "supply.name";
}
```

4.1.3 System Setup for Testing

The benchmarks of the sample applications were performed in Amazon EC2 ¹ instances. A distributed environment with configurations shown in Table 4.1 was setup on multiple EC2 instances.

4.1.4 Observations from DDE based applications

Some observations made during the development and the deployment of applications based on the DDE framework were:

- The program source code based on the DDE framework is short, as evident from the source code of the applications listed in section 4.1.

¹<http://aws.amazon.com/ec2>

- The worker isolate could not be run in browser because of current limitations of the Dart VM.
- Setting up the system for the first time was complicated because it consisted of several components that needed to be started up separately.
- Once the system was setup, adding workers and isolate systems to the nodes was easy. Thus, deploying an application to a distributed system was easy.
- Shutting down a single isolate or an isolate system did not effect other components and nodes.
- The decoupling of isolate systems, the MQS and the registry allowed each component to start up and shutdown without severely affecting other systems.
- The worker isolates developed for an application were loosely coupled. They were coupled only in terms of message.
- The REST API exposed by the registry simplified the deployment of isolate systems to different nodes.
- The Web interface of the Registry provided the ability to visualize and monitor the cluster.
- Messages sent to RabbitMQ were not lost during the restart of the system because they were persisted in RabbitMQ. The system continued to process messages after the restart.

4.2 Benchmarks

The benchmarks discussed here are the results of the evaluation of applications discussed in section 4.1. The results presented in this chapter were performed on Amazon EC2 servers² with configurations shown in table Table 4.1

4.2.1 Prefetch Count

Figure 4.1 shows the variation in overall message consumption when the number of consumers (separate isolate systems in separate nodes) were increased. Each line in the figure represents different values for the prefetch-count subsection 2.7.2 value set by the consumer while subscribing to the message broker system. Depending upon

²EC2 servers of data center in Ireland

Table 4.1: Specification of machines used for testing and benchmarking

Deployed Systems	Name	Specifications
RabbitMQ and a Message Queuing System	c3.2xLarge	8 core CPU, 28 ECU, 15 GiB Memory, SSD disk
Message Queuing System	m3.2xLarge	8 core CPU, 26 ECU, 30 GiB Memory, SSD disk
Registry and File Server	m3.xLarge	2 core CPU, 13 ECU, 15 GiB Memory, SSD disk
Isolate Systems (Nodes)	m3.xLarge	2 core CPU, 13 ECU, 15 GiB Memory, SSD disk

prefetch-count, increase, peak and decline of message consumption throughput can be seen for different numbers of consumers.

The increase in consumers had significant positive results for message throughput, but we can see that after around 8 consumers, adding more consumers had negative effect in the overall throughput.

Increasing prefetch count had more positive impact in message throughput compared to increasing consumers. For instance, the increase in consumers from 1 to 16 resulted in the rise of approximately 2000 messages per second, while the increase in prefetch count from 1 to 16 in the single consumer case resulted in an increase of message throughput by approximately 5500 messages per second.

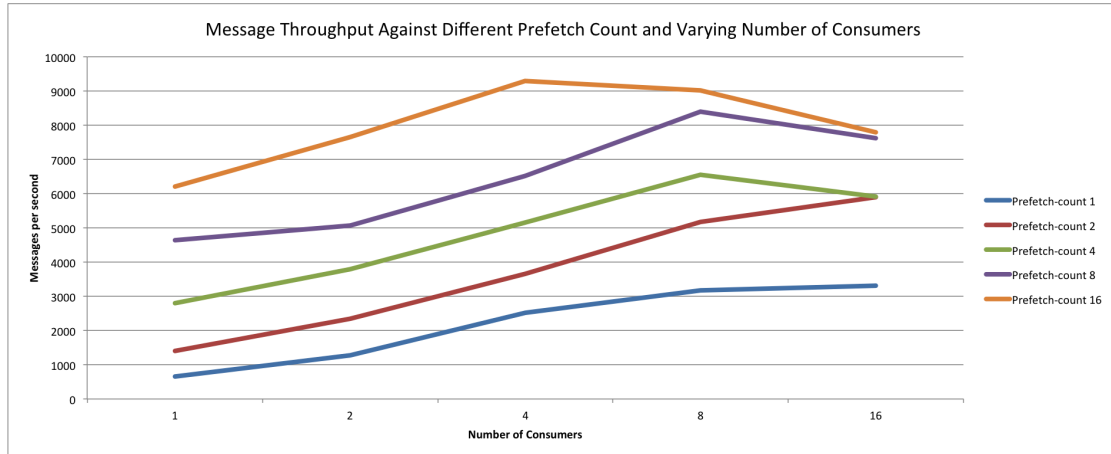


Figure 4.1: Message Throughput vs Number of consumers for varying prefetch-count (Higher is better)

4.2.2 Message Size

The result shown in Figure 4.2 was obtained by having eight consumers with a prefetch count of 8 to dequeue existing messages from a queue.

The negative effect on message consumption throughput of larger message is evident from the Figure 4.2. The decrease with message throughput was more prominent between message sizes of 256 bytes and 512 bytes than that of between 64 bytes and 256 bytes.

Similar observation can be made from Figure 4.3 which is a result of allowing a single worker isolate to create messages, except there was a slight rise in the throughput when the message size was increase from 512 bytes to 1024 bytes.




Figure 4.2: Message Consumption Throughput vs Message Size (Higher is better)




Figure 4.3: Message Production Throughput vs Message Size (Higher is better)

4.2.3 Number of Message Queuing Systems

Figure 4.4 is the result obtained by testing message consumptions for 1 to 32 consumers distributed to one to four MQS instances connected to the same message broker. Adding MQS instances clearly increased message throughput. Nevertheless, adding more than eight consumers per MQS instance had a negative effect, as we can see from the decline of throughput in the line charts for the one-instance case and the two-instance case. In the tests with one and two MQS instances, the optimum performance was seen when there were 8 consumers in total. But with four MQS instances, the throughput kept rising and supported up to 32 consumers without decline in the performance. Nevertheless, the rate of performance increase was not as much as the rate seen when scaling up from two MQS instances to four MQS instances.

The positive impact of scaling out MQS was seen not only on consumption throughput but also on production throughput. Figure 4.5 shows results of message production throughput by increasing the number of producers from one to eight, distributed across multiple MQS instances. The production rate measured here was the rate at which RabbitMQ enqueued the messages, not the rate at which a worker isolate produced messages.




Figure 4.4: Consumption Throughput on Scaled out Message Queuing System




Figure 4.5: Production Throughput on Scaling out Message Queuing System

4.2.4 Production Throughput of Isolates

In contrast to results of production throughput of messages as seen in Figure 4.3 and Figure 4.5, the result shown in Figure 4.6 measures the throughput of message production by the isolate system before it is sent to the MQS and broker for enqueueing. The production of a message in a Worker isolate was throttled, so that there is no immediate 'out of memory' error from an overwhelming production of messages. Throughput was throttled by delaying the production of messages by a random amount of time ranging from 20 - 500 microseconds. The observations made here are the average throughput of each worker isolate as well as average throughput of all the producer nodes connecting to a single MQS instance.

As observed in the Figure 4.6, the total production rate of messages sharply increased when increasing the number of producers. In contrast, the average production of a single node was lower with a higher number of producers.

In Dart version 1.7.2, time required to send a message of 'Map' datatype from one isolate to another consumed around 300 - 500 microseconds. But, when the same message was serialized to JSON String the time required to send a message dropped to 10 - 40 microseconds. It is an improvement in speed by one order of magnitude.




Figure 4.6: Production Throughput of Isolate

4.2.5 Simultaneous Production and Consumption

In contrast to previous benchmarks, which were measured with either only producers or only consumers, the benchmarks in Figure 4.7 and Figure 4.8 show the consumption rate and production rate of messages when they are run simultaneously but in different nodes. Compared to what was seen in Figure 4.1 and Figure 4.4, the consumption rate is lower in this case. Nevertheless, the production rate of messages remained almost equal to that which was observed in Figure 4.4 with a single MQS.

If we compare Figure 4.7 and Figure 4.8, scaling out with producers and consumers connecting to different MQS instances had a significant positive impact in overall message throughput. Especially, in this case the consumption rate increased by almost an order of magnitude than that of using a single MQS. A slight increase in production throughput was also observed with the additional MQS.

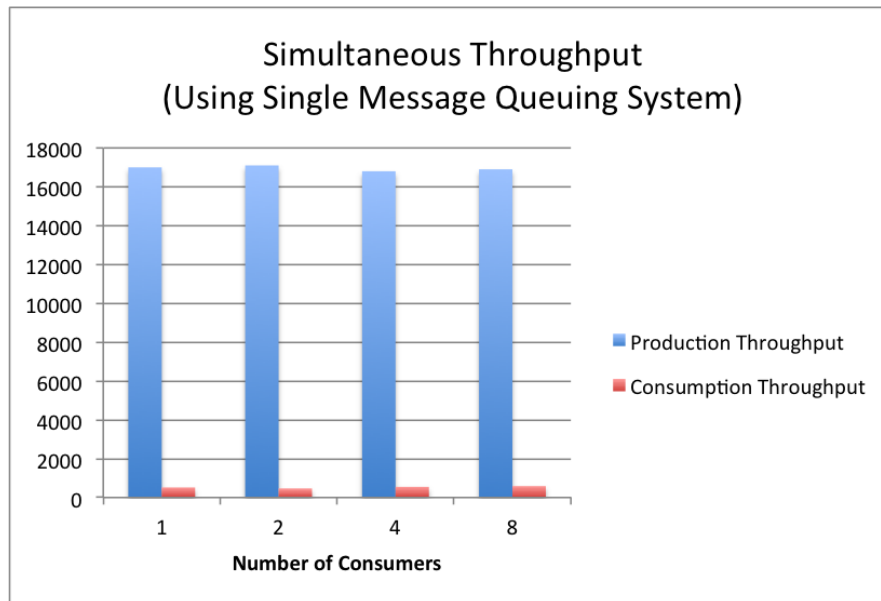


Figure 4.7: Throughput during simultaneous execution in a single MQS instance

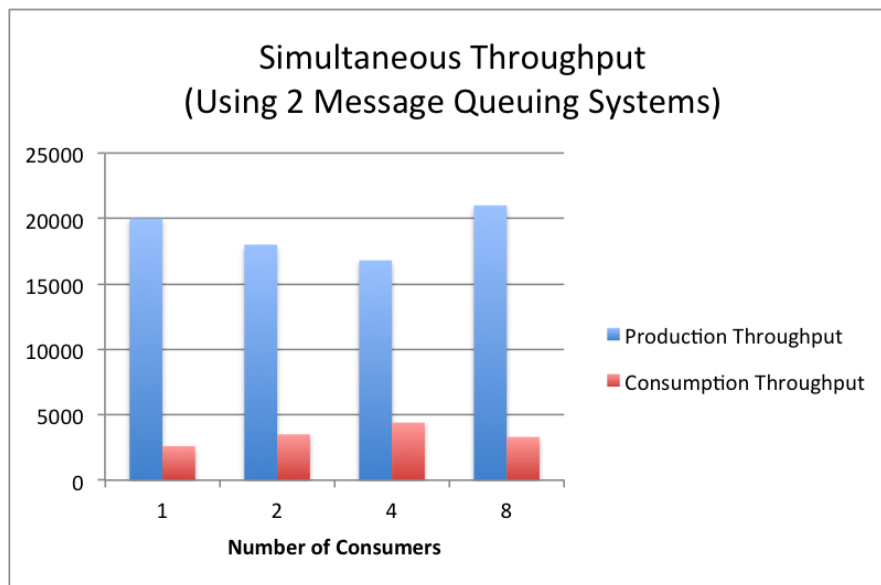


Figure 4.8: Throughput during simultaneous execution in two MQS instances

4.2.6 Throughput of Request-Reply

The request-reply application used for testing is different from the producer-consumer application in terms of how the messages are produced and consumed. In request-reply, a message is produced by the producer only when it receives a request from the requester (and the requester sends a request for another message only after it receives the reply) whereas, in producer-consumer application the producer creates messages regardless of the existence of consumers. Also, in the producer-consumer application any instance of a target worker isolate of any node may consume the message as it is designated only by the name of the worker isolate. But, in request-reply the reply message is sent to the worker isolate that originated the request.

For single execution of a message in producer-consumer, the steps required are enqueueing the message from producer and then dequeuing the message by the consumer. Whereas, in case of request-reply, first a request message must be enqueued and dequeued, then the reply message must be enqueued and dequeued. Thus, the steps required in request-reply are twice as much as in the producer-consumer case.

In Figure 4.9, when the number of consumers were increased we can see a linear increase in the production and consumption rate of message. However, adding more suppliers had very little effect in overall throughput. The maximum increase in throughput by adding suppliers was seen when there were more consumers.




Figure 4.9: Throughput of request-reply

4.2.7 Round Trip Time in Request-Reply

The results seen in Figure 4.10 is the benchmark of average round trip time (RTT) of messages in the request-reply case.

The result was obtained by attaching a timestamp in the message that was sent to the supplier; upon receiving the request message at the supplier, the timestamp was copied to the reply message and sent back to the sender(requester). The difference between the timestamp when the message was received at the requester and the timestamp contained in the message is the round trip time.

From the result of RTT shown in Figure 4.10, we can see that there were inconsistencies in RTT of messages with respect to number of suppliers as well as requesters. When there was only one supplier, the RTT of a message was usually higher compared to cases when there were more number of suppliers. Nevertheless, the least inconsistency was seen in case of 8 suppliers and the least latency was observed in case of single supplier with single consumer.




Figure 4.10: Round trip times of messages in request reply (Lower is better)

5 Discussions

5.1 Performance Findings

5.1.1 Effect of Prefetch Count

In the observations made in Figure 4.1, the throughput scaled almost linearly to 8 consumers for every prefetch-count (except when the prefetch-count was 16). Adding more consumers after 8 resulted in a decrease of throughput, which is quite similar to the result of benchmark [Figure 2.2] of RabbitMQ presented in subsection 2.7.2. In the benchmark of RabbitMQ, a decline of throughput was seen with 8 to 10 consumers.

The greater the number of consumers, the more work is required by RabbitMQ to keep track of them. Hence, it could be the same reason we saw the decline in performance with more consumers in Figure 4.1.

In a distributed system, a sudden performance change could be due to a number of factors, such as network speed, the system in which the isolate system is running or because of a bottleneck in the MQS. But each time the performance seems to drop after around 8 consumers. The number of consumers after which the decline in performance occurred in subsection 2.7.2 and Figure 4.1 suggests that the reason could be due to the broker handling too many consumers for a queue.

5.1.2 Effect of Message Size

From Figure 4.2 and Figure 4.3 from subsection 4.2.2, we observed the negative impact of increase in message size. Nevertheless, if we assess the amount by which the performance decreased, we can see that doubling the message size did not cause halving the performance. Thus, from this finding we can infer that instead of two separate chunks of messages, if we send messages as a big chunk the overall throughput increases. For instance, let's consider the results of 64-byte messages and 1024-byte messages. The throughput of 64-byte messages was 8,400 messages per second whereas that of 1024-byte was 3,097 messages per second. If we assume that the 1024-byte message was the concatenation of sixteen 64-byte messages, then the overall throughput would become 49,552 messages per second, which is an overall improvement of about 6 times. This means that for small messages, time is spent more on processing rather than on Input/Output.

The similar case for production throughput can be seen in Figure 4.3. With similar calculations for 64-byte messages and 1024-byte messages, if the messages are chunked together, we can gain approximately 13 times greater throughput.

Nevertheless, such concatenation may not always be feasible and may prove to be against the fair distribution of messages across systems. However, if it is appropriate to concat multiple messages in certain systems, then finding an optimum size (by considering the other requirements of the application) of a message which gives maximizes throughput might prove to be quite beneficial.

5.1.3 Effect of Scaling out the Message Queuing System

Scaling out the MQS clearly had positive impact on overall throughput. Higher number of MQS instances supported more consumers with good overall performance. The decline of performance seen in tests with 1 MQS instance and 2 MQS instances were probably because of the behavior of RabbitMQ when too many consumers are connected as we discussed in subsection 5.1.1. But, the test seen with four MQS instances suggests otherwise. With four MQS instances, 32 consumers were supported without a drop in performance. From which, we can speculate that the MQS might also be causing the limit in throughput, as the MQS has to distribute messages to the connected systems. As more consumers are connected to the MQS, more CPU time and memory are required.

5.1.4 Comparison of production throughputs

As seen in Figure 4.6 and Figure 4.3, the difference seen between the throughput at a worker isolate and throughput at RabbitMQ was quite significant. This is because the production of messages in an isolate were localized to one particular instance before the messages were transferred to the MQS over a WebSocket. Thus, it is natural the message production of isolates were high. Message production throughput increased with increasing the number of producers, but average throughput per isolate declined. One possible explanation for this behavior could be a bottleneck at the MQS to which the producers send messages. This explanation might seem strange because sending messages in an isolate are asynchronous, so it should not have been affected by an outside decoupled component. But, the reason behind this speculation is that since the MQS cannot accept messages at the rate they are produced, the messages get buffered in the internal queues of top level isolates. This takes up more heap memory reserved for the isolate causing slower execution. Similar effects probably occur in controller isolates, then in router isolates and ultimately in worker isolates. Hence, as the available heap memory of worker isolates becomes lower the production throughput decreases.

5.1.5 Effect of Simultaneous Production and Consumption

In Figure 4.7, we can see the negative impact of starting producers along with consumers when the producers were producing messages continuously. Again, when the producers and consumers were split to connect to separate instances of the MQS, the performance improvement was more prominent in message consumption.

By comparing Figure 4.7 and Figure 4.8, the culprit of this behavior seems to be the MQS. Even though the MQS has separate isolates for enqueueing and dequeuing, it has only one top level isolate subsection 3.2.3. Obviously when there is a large quantity of messages from a producer in queue, the dequeue requests that arrive from the consumer and the messages that arrive from the Dequeueer goes further back in the internal isolate queue. The improvement of consumption throughput seen in Figure 4.8 also supports this explanation because in this case, the MQS where consumers are connected does not have to deal with the large surge of production messages. Thus, the dequeue requests and dequeued messages from the dequeueer get processed much quicker compared to the previous case.

5.1.6 Throughput of Request-Reply

The benchmark of request-reply [subsection 4.2.6] showed almost linear improvement with increase in the number of requesters. But, increasing the number of suppliers from 1 to 8 had little effect in overall performance. The reason behind this could be the design of the program which uses the *ask* method to send messages. Based on how *ask* works as discussed in subsection 3.4.2, even a single supplier might not have been saturated with the number of requesters it could handle. A slight increase was seen with more suppliers, but this is quite minimal and could have been affected by other factors like RabbitMQ performance, change in network latency, etc. There is not enough data and evidence to support improvements in throughput with more suppliers. Testing with more consumers till the supplier saturates would yield a better understanding of what affects throughput in the request-reply case.

5.1.7 Round Trip Time of Request-Reply

The observation seen in Figure 4.10 was inconsistent. It is difficult to make any strong conclusions from the data. Nevertheless, it seems that the increase in number of suppliers decreases latency, as seen in case of eight suppliers against the single supplier.

There could be many factors affecting the result. Even a slight change in network latency can induce a significant change in time required to transfer a message.

5.2 The DDE Framework

According to the observations made after creating and testing the applications based on the framework (as discussed in subsection 4.1.4), we can infer the applications responded well when the system was scaled up. We can see a rise in message throughput when additional isolate systems are added. The rise in message throughput was even more when the number of Message Queuing Systems were increased.

During testing, restarting the message broker system (RabbitMQ) and restarting the message queuing systems did not result in failure of the whole system. The system operated normally as soon as they were back online. Thus, we can say that the system was highly available and it is suitable for systems that must be “run forever”.

5.3 Problems/Issues

5.3.1 Dart Induced Issues

Even though Dart is a fairly mature language, it still has some bugs and incomplete features. Some of the issues could be solved by workarounds while others could not be easily resolved. All the issues listed here were faced in the Dart version 1.7.2.

Dart’s documentation for isolate [Goo] mentions that there is experimental support for following methods in isolate: *addErrorListener*, *addOnExitListener*, *kill*, *pause*, *ping*, *removeErrorListener*, *removeOnExitListener*, *resume*. But, while developing the framework, it was found that these methods were not implemented. The implementation of these functionalities would have particularly helped in the implementation of the supervision strategy. Moreover, the implementation of features like ‘hot deployment’, ‘migration’ and isolate termination would have been cleaner and simpler compared to the current workarounds present in the DDE framework.

We have already discussed the workaround implementations for *kill* in subsection 3.4.2 and *ping* in section 3.2.1.

Another issue found in Dart VM for browsers which made it impossible for an isolate system designed in DDE framework to run was the lack of support for multi-level isolates. The ‘Dartium’ browser ¹, only supports spawning of isolates up to one level deep. i.e. Spawning of an isolate by the spawned isolate was not allowed. The reason behind the lack of this implementation is not clear.

¹a variant of Chromium browser with built in dart VM

5.3.2 STOMP Library for Dart

As mentioned in subsection 2.8.2, a third-party library for the STOMP client in Dart was used from Dart's 'pub'. Although it had good support for STOMP 1.2, it was missing a feature to set the prefetch count [subsection 2.7.2] while subscribing as a consumer. Without setting the prefetch count the system became unusable, especially when there were many messages in the queue.

As the project was hosted in git hub ² and is licensed under Apache 2.0 ³, I forked ⁴ the repository and added the functionality so that the STOMP client could support setting the prefetch-count. After making modifications and testing, I created a pull-request ⁵, which was swiftly merged ⁶ by the original developer and was updated in Dart's pub as a new version.

²<http://github.com>

³<https://github.com/rikulo/stomp/blob/master/LICENSE>

⁴make copy of the source code and add custom modifications

⁵request to merge my changes into the master branch of the original repository

⁶<https://github.com/rikulo/stomp/pull/15>

6 Conclusion

The DDE framework, created as a result of this thesis, shows that isolates of the Dart language do not have to be limited to a single virtual machine.

The DDE framework provides structure for developers to create applications based on the actor programming model in the Dart language. This framework uses the existing actor-like nature of isolates and provides easy to use functions to make message sending similar to the actor programming model.

As a result of using this framework, applications become easily scalable and offer higher availability with virtually no down time during code updates. Scaling up can be performed in several ways: deploying more isolates in an isolate system, replicating isolate systems across multiple machines, increasing the number of message queuing systems, forming a cluster of message broker systems (RabbitMQ) or the combination of these.

Moreover, the DDE framework offers developers the capabilities to monitor code updates and restart worker isolates automatically. Moreover, the web interface and REST API available in the 'Registry' allow developers to visualize, easily deploy, terminate individual or groups of isolates, and even shutdown isolate systems.

The message broker system (RabbitMQ) of the DDE framework allows applications to be loosely coupled and improves the overall fault-tolerance of the system. Moreover, the clustering ability and message persistence provided by RabbitMQ allows the messages to be saved in several machines ensuring message durability.

Nevertheless, the DDE framework has not yet implemented some features required by certain users. For example, it does not have a security implementation of its own.

The preliminary tests and benchmarks of the DDE framework are encouraging. For the proof of concept implementation, the throughput of message production and consumption were good. With further profiling and optimization throughput and latency improvements can be achieved.

7 Future Directions

The work done in the DDE framework in this thesis may be extended and improved in number of directions. First, the most important missing feature, security, requires a concept and implementation. The DDE framework heavily relies on the underlying transport layer to provide security, additional application level security would provide the DDE framework more safety.

Another important feature is to support developers in testability of DDE applications. This would improve programmability and help developers to write quality codes.

The implementation of supervision strategy is far from perfect in the presented DDE framework. The current implementation of the supervision strategy does not properly follow the “Let it Crash” [subsection 2.3.1] philosophy. As discussed in subsection 5.3.1, the implementation of this feature would have been simpler and cleaner had Dart supported all of the unimplemented features of isolates. Nevertheless, the Dart language is evolving quickly and those features may eventually be implemented. However, even if those features do not make into the major version of dart any time soon, it is possible to make implementations using workarounds.

Another interesting idea would be the introduction of adaptive load balancing and automatic migration of isolates based on different heuristics like load on a system, data locality or a user’s geographic location.

One feature to improve the fault-tolerance of the system would be to change the implementation for FileMonitor [section 3.2.1]. The current implementation of the FileMonitor simply monitors the file from a remote location and periodically fetches it to calculate the checksum. In case of modification in the file, the FileMonitor informs the controller which instructs a router to respawn the worker regardless of the contents of the source. This means there is the vulnerability to deploy incorrect code which could result in failing to start up the isolate. The ability to rollback to a previous working version of the source in case of failure of spawning would improve the resilience of the framework.

The ideas discussed in these chapters are only a few of many other improvements that can be made in the framework.

Acronyms

JSON JavaScript Object Notation.

MBS Message Broker System.

MQS Message Queuing System.

URI Uniform Resource Identifier.

UUID Universally Unique Identifier.

List of Figures

2.1	Hierarchy in actor system	10
2.2	Chart showing performance variation when prefetch count is changed in the consumer	21
3.1	Architecture of the framework	25
3.2	Adding a Worker isolate to an isolate system	27
3.3	Proxy Worker	32
3.4	Enqueuing a message	40
3.5	Dequeuing a message	41
4.1	Message Throughput vs Number of consumers for varying prefetch-count	53
4.2	Message Consumption Throughput vs Message Size	54
4.3	Message Production Throughput vs Message Size	55
4.4	Consumption Throughput on Scaled out Message Queuing System . . .	56
4.5	Production Throughput on Scaling out Message Queuing System . . .	56
4.6	Production Throughput of Isolate	58
4.7	Throughput during simultaneous execution in single MQS	59
4.8	Throughput during simultaneous execution in two MQS	59
4.9	Throughput of request-reply	60
4.10	Round trip times of messages in request reply	61

List of Tables

3.1	Routing techniques provided by the framework	30
3.2	Endpoints exposed by the registry	34
3.3	Dependent libraries of the framework	45
4.1	Specification of machines used for testing and benchmarking	53

Bibliography

- [Agh85] G. A. Agha. *ACTORS: A Model of Concurrent Computation in Distributed Systems*. English. 1985, pp. 34–35.
- [AH85] G. Agha and C. Hewitt. “Concurrent programming using actors: Exploiting large-scale parallelism.” English. In: *Foundations of Software Technology and Theoretical Computer Science*. Ed. by S. Maheshwari. Vol. 206. Lecture Notes in Computer Science. Springer Berlin Heidelberg, 1985, pp. 19–41. ISBN: 978-3-540-16042-7. DOI: 10.1007/3-540-16042-6_2.
- [Arm07] J. Armstrong. “A History of Erlang.” In: *Proceedings of the Third ACM SIGPLAN Conference on History of Programming Languages*. HOPL III. San Diego, California: ACM, 2007, pages. ISBN: 978-1-59593-766-7. DOI: 10.1145/1238844.1238850.
- [Arm10] J. Armstrong. “Erlang.” In: *Commun. ACM* 53.9 (Sept. 2010), pp. 68–75. ISSN: 0001-0782. DOI: 10.1145/1810891.1810910.
- [Cel11] J. Celko. *Joe Celko’s SQL for Smarties: Advanced SQL Programming*. Morgan Kaufmann, 2011. ISBN: 978-0-12-382022-8.
- [ECM14] ECMA. *ECMA-408: Dart Programming Language Specification*. 2014. Geneva, Switzerland: ECMA (European Association for Standardizing Information and Communication Systems), June 2014.
- [Erb12] B. Erb. “Concurrent Programming for Scalable Web Architectures.” Diploma Thesis. Institute of Distributed Systems, Ulm University, Apr. 2012.
- [FM11] I. Fette and A. Melnikov. *The WebSocket Protocol*. RFC 6455 (Proposed Standard). Internet Engineering Task Force, Dec. 2011.
- [Goo] Google. *Dart API Reference*. <https://api.dartlang.org/apidocs/channels/stable/dartdoc-viewer/dart:isolate>. Isolate. Last Accessed: 2014-11-19.
- [Hal] P. Haller. *Scala Actors*. <http://www.scala-lang.org/old/node/242>. Last Accessed: 2014-11-19.

- [HO09] P. Haller and M. Odersky. "Scala Actors: Unifying Thread-based and Event-based Programming." In: *Theor. Comput. Sci.* 410.2-3 (Feb. 2009), pp. 202–220. ISSN: 0304-3975. DOI: 10.1016/j.tcs.2008.09.019.
- [HT] P. Haller and S. Tu. *Scala Actors API*. <http://docs.scala-lang.org/overviews/core/actors.html>. Last Accessed: 2014-11-30.
- [Inca] T. Inc. *Actor Systems*. <http://doc.akka.io/docs/akka/snapshot/general/actor-systems.html>. Last Accessed: 2014-12-01.
- [Incb] T. Inc. *Akka Toolkit*. <http://akka.io>. Last Accessed: 2014-11-19.
- [Incc] T. Inc. *Akka Toolkit*. <http://doc.akka.io/docs/akka/2.3.7/AkkaJava.pdf>. Last Accessed: 2014-11-19.
- [Kat12] S. L. Kathy Walrath. *Dart: Up and Running*. Last Accessed: 2014-11-19. O'Reilly Media, Oct. 2012. ISBN: 978-1-4493-3084-2.
- [Lad] S. Ladd. *Dart Is Not the Language You Think It Is*. <http://radar.oreilly.com/2013/05/dart-is-not-the-language-you-think-it-is.html>. Last Accessed: 2014-12-01.
- [Nin] C. Ninnars. *The Actor Model (everything you wanted to know, but were afraid to ask)*.
- [Nor] P. Nordwall. *Running a 2400 Akka Nodes Cluster on Google Compute Engine*. <http://typesafe.com/blog/running-a-2400-akka-nodes-cluster-on-google-compute-engine>. Last Accessed: 2014-12-01.
- [OR14] M. Odersky and T. Rompf. "Unifying Functional and Object-oriented Programming with Scala." In: *Commun. ACM* 57.4 (Apr. 2014), pp. 76–86. ISSN: 0001-0782. DOI: 10.1145/2591013.
- [Pip] A. Piper. *Choosing Your Messaging Protocol: AMQP, MQTT, or STOMP*. <http://blogs.vmware.com/vfabric/2013/02/choosing-your-messaging-protocol-amqp-mqtt-or-stomp.html>. Last Accessed: 2014-11-22.
- [Sofa] U. o. B. Software Languages Lab. *Concurrent Programming with Actors*. <http://soft.vub.ac.be/amop/at/tutorial/actors>. Last Accessed: 2014-11-30.
- [Sofb] P. Software. *RabbitMQ Performance Measurements, part 2*. <http://www.rabbitmq.com/getstarted.html>. Last Accessed: 2014-12-02.
- [Sofc] P. Software. *RabbitMQ Performance Measurements, part 2*. <http://www.rabbitmq.com/blog/2012/04/25/rabbitmq-performance-measurements-part-2/>. Last Accessed: 2014-11-20.
- [Sofd] P. Software. *Sizing your Rabbits*. <https://www.rabbitmq.com/blog/2011/09/24/sizing-your-rabbits/>. Last Accessed: 2014-11-22.

Bibliography

- [Sofe] P. Software. *What can RabbitMQ do for you?* <http://www.rabbitmq.com/features.html>. Last Accessed: 2014-11-20.
- [Soff] P. Software. *Which protocols does RabbitMQ support?* <http://www.rabbitmq.com/protocols.html>. Last Accessed: 2014-11-20.
- [STO] STOMP. *STOMP - Simple Text Oriented Messaging Protocol*. <http://stomp.github.io/stomp-specification-1.2.html>. Last Accessed: 2014-11-20.
- [Vin07] S. Vinoski. "Concurrency with Erlang." In: *Internet Computing, IEEE* 11.5 (Sept. 2007), pp. 90–93. ISSN: 1089-7801. DOI: 10.1109/MIC.2007.104.